



## Operations Research

Publication details, including instructions for authors and subscription information:  
<http://pubsonline.informs.org>

### Multistage Robust Mixed-Integer Optimization with Adaptive Partitions

Dimitris Bertsimas, Iain Dunning

To cite this article:

Dimitris Bertsimas, Iain Dunning (2016) Multistage Robust Mixed-Integer Optimization with Adaptive Partitions. Operations Research 64(4):980-998. <https://doi.org/10.1287/opre.2016.1515>

Full terms and conditions of use: <https://pubsonline.informs.org/page/terms-and-conditions>

This article may be used only for the purposes of research, teaching, and/or private study. Commercial use or systematic downloading (by robots or other automatic processes) is prohibited without explicit Publisher approval, unless otherwise noted. For more information, contact [permissions@informs.org](mailto:permissions@informs.org).

The Publisher does not warrant or guarantee the article's accuracy, completeness, merchantability, fitness for a particular purpose, or non-infringement. Descriptions of, or references to, products or publications, or inclusion of an advertisement in this article, neither constitutes nor implies a guarantee, endorsement, or support of claims made of that product, publication, or service.

Copyright © 2016, INFORMS

Please scroll down for article—it is on subsequent pages

INFORMS is the largest professional society in the world for professionals in the fields of operations research, management science, and analytics.

For more information on INFORMS, its publications, membership, or meetings visit <http://www.informs.org>

# Multistage Robust Mixed-Integer Optimization with Adaptive Partitions

Dimitris Bertsimas, Iain Dunning

Operations Research Center and Sloan School of Management, Massachusetts Institute of Technology, Cambridge, Massachusetts 02139  
{[dbertsim@mit.edu](mailto:dbertsim@mit.edu), [idunning@mit.edu](mailto:idunning@mit.edu)}

We present a new partition-and-bound method for multistage adaptive mixed-integer optimization (AMIO) problems that extends previous work on finite adaptability. The approach analyzes the optimal solution to a static (nonadaptive) version of an AMIO problem to gain insight into which regions of the uncertainty set are restricting the objective function value. We use this information to construct partitions in the uncertainty set, leading to a finitely adaptable formulation of the problem. We use the same information to determine a lower bound on the fully adaptive solution. The method repeats this process iteratively to further improve the objective until a desired gap is reached. We provide theoretical motivation for this method, and characterize its convergence properties and the growth in the number of partitions. Using these insights, we propose and evaluate enhancements to the method such as warm starts and smarter partition creation. We describe in detail how to apply finite adaptability to multistage AMIO problems to appropriately address nonanticipativity restrictions. Finally, we demonstrate in computational experiments that the method can provide substantial improvements over a nonadaptive solution and existing methods for problems described in the literature. In particular, we find that our method produces high-quality solutions versus the amount of computational effort, even as the problem scales in the number of time stages and the number of decision variables.

**Keywords:** adaptive optimization; robust optimization; integer optimization.

**Subject classifications:** programming: stochastic, linear, integer.

**Area of review:** Optimization.

**History:** Received November 2014; revisions received August 2015, January 2016; accepted April 2016. Published online in *Articles in Advance* June 24, 2016.

## 1. Introduction

Robust optimization (RO) is a methodology for addressing uncertainty in optimization problems where *uncertain parameters* are modeled as belonging to *uncertainty sets*, which contrasts with the more traditional approach of modeling uncertainty with probability distributions. We solve a RO problem by optimizing with respect to the worst-case realization of the uncertain parameters over the uncertainty set  $\Xi$ ,

$$\begin{aligned} \max_{\mathbf{x} \in \mathcal{X}} \min_{\xi \in \Xi} c(\xi, \mathbf{x}) \\ \text{subject to } \mathbf{g}(\xi, \mathbf{x}) \leq \mathbf{0} \quad \forall \xi \in \Xi, \end{aligned} \quad (1)$$

where  $\mathbf{x}$  is the vector of decision variables and  $\xi$  is the vector representing the uncertain parameters. For a survey of RO, we refer the reader to Bertsimas et al. (2011a) and Ben-Tal et al. (2009). The optimization problem (1) commonly has an infinite number of constraints but is tractable for many uncertainty sets of interest by reformulating the problem to obtain a deterministic optimization problem of finite size, or by generating constraints only as needed until the solution is optimal and feasible with respect to the uncertain constraints (e.g., Fischetti and Monaci 2012, Bertsimas et al. 2016).

RO was introduced in the context of single-stage problems as a way to address uncertainty in problem data. Since

then, it has been shown to be a more general methodology to address multistage optimization problems, where the uncertain parameters are revealed over time. In most multistage problems, our future *recourse* decisions can be made with knowledge of at least some of the uncertain parameters: those realized by the time those decisions must be made. If we place no restrictions on the functional relationship between uncertain parameters and recourse decisions, which is the ideal case, then the resulting problem is known to be hard in the computational complexity sense (Ben-Tal et al. 2004), and has proved to be hard to solve in practice. Some work on addressing this fully adaptive case includes a Benders' decomposition-like approach to solve a unit commitment problem in Bertsimas et al. (2013), and the generalized "column-and-constraint generation" framework for two-stage-continuous problems of Zeng and Zhao (2013). A goal of the adaptive optimization (AO) literature then is to provide methods to approximate and approach this fully adaptable ideal and to quantify, if possible, how good these approximations are.

Perhaps the most popular approach in the AO literature to date has been to restrict the recourse decisions to be simple functions of the uncertain parameters. Affine adaptability, also known as linear decision rules, was introduced to the RO setting in Ben-Tal et al. (2004) and has proven

to be useful in a wide variety of settings, including control problems (Goulart et al. 2006), supply chain problems (Ben-Tal et al. 2005), and project management (Chen et al. 2007). The choice of these linear decision rules is primarily one of computational practicality with no real expectation that they will be optimal in general (where optimality would be finding an affinely adaptive policy that performs as well as the fully adaptive policy). Many of the papers cited here demonstrate these shortcomings with both examples and more realistic experiments, and many researchers have characterized the conditions under which, and to what degree, affine adaptability performs well, e.g., Bertsimas et al. (2010), Bertsimas and Goyal (2012), Bertsimas and Bidkhori (2015). More complicated approaches trade computational efficiency for optimality, for example, by using polynomial adaptability (Bertsimas et al. 2011b) and deflected linear decision rules (Chen et al. 2008). A major shortcoming of all these approaches is that they do not allow for discrete recourse decisions, significantly reducing their modeling power for adaptive mixed-integer optimization (AMIO) problems. Bertsimas and Caramanis (2007) acknowledge this and propose a decision rule for integer decisions as part of their sampling approach. Bertsimas and Georghiou (2015) propose a cutting-plane method that allows for piecewise linear decision rules for continuous recourse decisions and piecewise constant decision rules for integer recourse decisions. This approach was shown to be able to find good solutions to small problems, but does not scale well to problems with many time stages. A recent extension by Bertsimas and Georghiou (2014) for binary recourse decisions that allows for a reformulation demonstrates substantially improved performance on multistage problems.

An alternative approach is finite adaptability, where instead of focusing on the functional form of the recourse decisions, we partition the uncertainty set and have different recourse decisions for each partition. This can be viewed as modeling the recourse decisions as piecewise constant functions of the uncertain parameters, with the domain of each piece (or equivalently the uncertainty set partitioning) either fixed, or decided endogenously as part of the optimization problem. One of the key benefits of this approach is that it handles discrete recourse decisions naturally, suggesting it is a good choice for AMIO. The quality of solutions obtained with a finite adaptability approach rests entirely on how the partitions are selected. For example, in Vayanos et al. (2011), the uncertainty set is partitioned ahead of time using hyperrectangles. A bilinear optimization problem that decides the best two-partition recourse decision is proposed in Bertsimas and Caramanis (2010), and the gap between a fully adaptive solution and a finitely adaptive solution is bounded. Finally, Hanasusanto et al. (2014) propose a general methodology for obtaining optimal  $K$ -partition recourse decisions for two-stage binary optimization problems, and characterize when their approach is able to match the fully adaptive solution. They

are able to solve a variety of problems, although even problems with as few as three partitions and 20 binary second-stage variables cannot be solved to optimality within two hours. While this recent work represents great progress in AMIO, the size of AMIO problems that can be solved in a reasonable amount of time remains small and lags significantly behind the size of tractable continuous AO problems, particularly in the multistage case.

In this paper, we will consider multistage AMIO problems of the general form

$$\begin{aligned} z_{\text{full}} = \min_{\mathbf{x} \in \mathcal{X}} \max_{\xi \in \Xi} \sum_{t=1}^T \mathbf{c}'(\xi) \cdot \mathbf{x}'(\xi^1, \dots, \xi^{t-1}) \\ \text{subject to} \quad \sum_{t=1}^T \mathbf{A}'(\xi) \cdot \mathbf{x}'(\xi^1, \dots, \xi^{t-1}) \leq \mathbf{b}(\xi) \quad (2) \\ \forall \xi = (\xi^1, \dots, \xi^T) \in \Xi \end{aligned}$$

where  $\xi$  represents the uncertain parameters,  $\mathbf{A}$ ,  $\mathbf{b}$ , and  $\mathbf{c}$  define linear constraints and the objective, respectively, (all themselves affine functions of  $\xi$ ), and  $\mathbf{x}'$  is the vector of decision variables for stage  $t$ , which can be a mixture of discrete and continuous variables as captured by the deterministic set of constraints  $\mathcal{X}$ . For  $t = 1$ , we note that  $\mathbf{x}^1$  is a nonadaptive here-and-now decision that does not depend on  $\xi$ .

The key contribution of this paper is a new “partition-and-bound” method for solving the multistage AMIO problem (2) that scales better to larger instances than alternative methods. In particular, as our method and the ones that we are comparing ourselves with do not necessarily find the fully adaptive solution, the metric by which we assess a method is the quality of the solutions it finds versus the computational effort required to obtain these solutions. The method builds on the finite adaptability approach to AMIO by using the information obtained about the (possibly already partitioned) uncertainty set when solving an AO problem to select a new partitioning scheme. The same information is used to calculate a lower bound on the fully adaptive solution by adapting the sampling proposal of Hadjiyiannis et al. (2011). This process of solving and updating the partitioning can be repeated until a termination criterion in terms of the bound gap, or some other limit (total iterations or computation time) is reached. Our proposal can be extended in many ways: in this paper, we show that affine adaptability for continuous recourse decisions can be trivially incorporated into our approach (leading to piecewise affine decision rules), but other extensions such as quadratic objectives and constraints can be also easily incorporated. We also give a detailed treatment to the interactions between finite adaptability and the need to enforce *nonanticipativity* for multistage problems—the natural notion that a decision at time  $t$  must be made without certain knowledge of uncertain parameters that will be revealed after time  $t$ .

Our approach lies somewhere between those that determine the optimal partitions simultaneously with the values of the decisions (Bertsimas and Caramanis 2010, Hanasusanto et al. 2014) and methods that select a partitioning a priori (Vayanos et al. 2011). While finalizing this paper, we became aware of a very recently submitted paper by Postek and Den Hertog (2014) that presents an iterative finite partitioning approach similar to the idea we present in this paper. Throughout this paper we will contrast and compare our approach to this alternative, from a theoretical and computational point of view.

After the initial submission of this paper, Song and Luedtke (2014) proposed an “adaptive partition-based approach” for two-stage stochastic programs (SP) that partitions sets of scenarios. Their approach uses the solutions to these relaxed problems to change the partitioning scheme—a scenario-based analogue to the approach of this paper. An interesting point of difference is the perspective of the two approaches: in the SP case, we start with many individual scenarios that we group into partitions, whereas in the RO approach, we start with a single set and seek to split apart into partitions. Song and Luedtke (2014) build on other similar aggregation techniques discussed in the SP literature, especially work aimed at improving the tractability of sample average approximation. We refer interested readers to their work for a more comprehensive treatment of the related SP literature.

We have structured the paper as follows:

- In Section 2, we present a theoretical result describing a property that a new partitioning of an uncertainty set must have to improve a solution of an AMIO over an existing partitioning scheme. We use this result to propose a method of constructing partitions using the solution of AMIO problems that is inspired by Voronoi diagrams. In this section, we initially focus on the two-stage case: By combining this with a lower bounding method, we obtain an iterative partition-and-bound method for two-stage AMIO. We explain how affine adaptability can be incorporated into the method, and provide a worked example to aid in implementation.

- In Section 3, we analyze our method from two points of view. We show that our method has a desirable nonincreasing property that enables computational “warm starts,” but provide an example that demonstrates that there is no general guarantee of convergence for the solutions produced by our method to the fully adaptive solution. We also consider the question of the growth in the number of partitions with increasing partitions. We show that while the number of partitions can grow very large after many iterations, the quality of the solution those partitions lead to for a fixed computational budget is the more relevant quantity. We also present a small modification to the method that can greatly reduce the number of partitions necessary.

- In Section 4, we generalize the method to the multistage case. We first demonstrate the issues that must be

considered in choosing the partitioning scheme while satisfying nonanticipativity. We then describe the general multistage partitioning scheme and a simple routine to determine the minimum set of nonanticipativity constraints to add.

- In Section 5, we provide results of computational experiments that show that our method produces good solutions versus total computation time. In particular, we solve instances of a capital budgeting problem and a multistage lot-sizing problem, including a comparison with the methods of Hanasusanto et al. (2014), Bertsimas and Georghiou (2015), and Postek and Den Hertog (2014).

- In the final Section 6, we provide some concluding remarks and thoughts on future directions.

## Notation

We use lowercase, nonbold face symbols for scalar quantities (e.g.,  $z \in \mathbb{R}$ ) and a bold face for vectors (e.g.,  $\mathbf{x} \in \mathbb{R}^n$ ). Matrices are denoted by bold face uppercase symbols, e.g.,  $\mathbf{A} \in \mathbb{R}^{m \times n}$ . We use  $\Xi$  to denote an uncertainty set and  $\xi$  to denote uncertain parameters. Our uncertain parameters are typically a vector of vectors, i.e.,

$$\xi = (\xi^1, \dots, \xi^t, \dots, \xi^T), \quad (3)$$

where  $\xi^t$  is the vector of uncertain parameters for stage  $t$ . The  $j$ th component of the uncertain parameters for time stage  $t$  is expressed with a subscript, e.g.,  $\xi_j^t$ . For particular realizations of uncertain parameters, we use a hat. For example,  $\hat{\xi}_i$  is the  $i$ th realization and  $\hat{\xi}_i^t$  is the vector of uncertain parameters for sample  $i$  and time stage  $t$ .

## 2. Partition-and-Bound Method for Two-Stage AMIO

In this section, we describe a partition-and-bound method that takes a finite adaptability approach to two-stage AMIO problems. In particular, we consider the following problem:

$$\begin{aligned} \min_{\mathbf{x} \in \mathcal{X}, z} \quad & z \\ \text{subject to} \quad & \mathbf{c}^1(\xi) \cdot \mathbf{x}^1 + \mathbf{c}^2(\xi) \cdot \mathbf{x}^2(\xi) \leq z \quad \forall \xi \in \Xi \\ & \mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\xi) \cdot \mathbf{x}^2(\xi) \leq b_i(\xi) \\ & \forall \xi \in \Xi, i \in \mathcal{I}, \end{aligned} \quad (4)$$

which has  $m$  linear uncertain constraints ( $\mathcal{I} = \{1, \dots, m\}$ ) and the objective expressed in epigraph form. This is a special case of the multistage problem (2), where  $T = 2$ . We first demonstrate some properties of AO solutions that motivate the design of our proposed method, then discuss the implementation details.

### 2.1. Partitioning Scheme

Finite adaptability is a simple approach for approximating fully adaptive wait-and-see decisions that can vary with the uncertain parameters. In the context of our two-stage problem (4), we can view finite adaptability as a restriction of



the decisions  $\mathbf{x}^2(\xi)$  to the class of piecewise constant policies, i.e.,

$$\mathbf{x}^2(\xi) = \begin{cases} \mathbf{x}_1^2, & \forall \xi \in \Xi_1, \\ \vdots \\ \mathbf{x}_K^2, & \forall \xi \in \Xi_K, \end{cases} \quad (5)$$

where  $\Xi_k$ ,  $k \in \{1, \dots, K\}$  defines a partitioning of  $\Xi$  and  $\mathbf{x}_k^2$  is the implemented decision if the realized value of  $\xi$  is in the partition  $\Xi_k$ . If the realized value lies on the boundary of two partitions, then we may select either decision, as the decision associated with each partition is feasible for all values of the uncertain parameters on the boundary. We will converge in the limit to the fully adaptive solution as the number of partitions grows and the diameter of the largest partition approaches 0. Therefore, there is a natural trade-off between the number of partitions (and the solution time) and how “close” our solution is to the fully adaptive solution, and a goal when using finite adaptability is to identify a partitioning scheme that is near the efficient frontier of this trade-off.

To motivate our selection of partitioning scheme, we will consider a variant of (4) where the second-stage decisions cannot change with  $\xi$ , i.e., a static policy:

$$\begin{aligned} z_{\text{static}} = \min_{\mathbf{x} \in \mathcal{X}, z} \quad & \\ \text{subject to} \quad & \mathbf{c}^1 \cdot \mathbf{x}^1 + \mathbf{c}^2 \cdot \mathbf{x}^2 \leq z \\ & \mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\xi) \cdot \mathbf{x}^2 \leq b_i(\xi) \\ & \forall \xi \in \Xi, i \in \mathcal{I}. \end{aligned} \quad (6)$$

We will initially restrict ourselves to the case where all decision variables are continuous, the uncertainty set is polyhedral (which ensures that a cutting plane method terminates), the objective coefficients are certain, and the problem is feasible and bounded. Consider using a cutting-plane method to solve this static problem, that is, we iteratively solve a deterministic relaxation of the RO problem with a finite subset of the possible constraints and add violated constraints from the full RO problem until the solution of the relaxation is feasible with respect to all uncertain parameters in the uncertainty set. For each constraint  $i$ , we define  $\mathcal{A}_i$  to be a set of *active uncertain parameters*  $\hat{\xi}$  corresponding to the generated constraints (cuts) that have zero slack. These sets may be empty, singletons, or even have multiple elements (multiple cuts for  $i$  with zero slack). After solving for the static policy, we consider partitioning the uncertainty set into two sets,  $\Xi_1$  and  $\Xi_2$ , that is,

$$\begin{aligned} z_{\text{part}} = \min_{\mathbf{x} \in \mathcal{X}, z} \quad & \\ \text{subject to} \quad & \mathbf{c}^1 \cdot \mathbf{x}^1 + \mathbf{c}^2 \cdot \mathbf{x}_j^2 \leq z \quad \forall j \in \{1, 2\} \\ & \mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\xi) \cdot \mathbf{x}_j^2 \leq b_i(\xi) \\ & \forall \xi \in \Xi_j, j \in \{1, 2\}, i \in \mathcal{I}, \end{aligned} \quad (7)$$

which we hope will have a solution that is better than the static policy (that is,  $z_{\text{part}} < z_{\text{static}}$ ). Let  $\mathcal{A} = \bigcup_i \mathcal{A}_i$ . We now seek to show that whether this improvement occurs or not depends on the nature of the partitions  $\Xi_1$  and  $\Xi_2$ .

**THEOREM 1.** *If  $\mathcal{A}$ , the set of all active uncertain parameters for the static problem (6), satisfies either  $\mathcal{A} \subset \Xi_1$  or  $\mathcal{A} \subset \Xi_2$ , then  $z_{\text{part}} = z_{\text{static}}$ . Otherwise,  $z_{\text{part}} \leq z_{\text{static}}$ .*


**PROOF OF THEOREM 1.** Suppose that the set of active uncertain parameters  $\mathcal{A}$  is a subset of  $\Xi_1$ . Consider solving the problem

$$\begin{aligned} z_{\text{half}} = \min_{\mathbf{x} \in \mathcal{X}, z} \quad & \\ \text{subject to} \quad & \mathbf{c}^1 \cdot \mathbf{x}^1 + \mathbf{c}^2 \cdot \mathbf{x}_1^2 \leq z \\ & \mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\xi) \cdot \mathbf{x}_1^2 \leq b_i(\xi) \\ & \forall \xi \in \Xi_1, i \in \mathcal{I}, \end{aligned} \quad (8)$$

which can be viewed as relaxation of (6) as we must only satisfy the constraints for a subset of  $\Xi$ . Despite being a relaxation, we still have  $z_{\text{half}} = z_{\text{static}}$ , as every cut with minimal slack in (6) will still be a valid cut for (8), and thus the solution will be equal as we have not relaxed any of the constraints that bound the previous solution. Note that this does not preclude the possibility that the equivalent problem to (8) for the other partition  $\Xi_2$  has an objective function value better than  $z_{\text{static}}$ , but as  $z_{\text{part}}$  is taken to be the maximum of the objective function values corresponding to the two partitions in (7), there can be no improvement of objective function value overall (i.e.,  $z_{\text{half}} = z_{\text{part}} = z_{\text{static}}$ ). On the other hand, if we have partitions such that  $\mathcal{A} \not\subset \Xi_1$  and  $\mathcal{A} \not\subset \Xi_2$ , then the possibility of improvement exists: at least one zero slack cut will not be valid for  $z_{\text{half}}$ , so the solution to (8) is not restricted to having the same solution as (6).  $\square$

This result naturally leads us to consider partitioning schemes that will achieve this required condition, and thus enable improvement. A partitioning scheme that achieves this is one in which every  $\hat{\xi} \in \mathcal{A}$  lies in its own partition. A simple way (conceptually and computationally) to construct such a set of partitions is to use a *Voronoi diagram* (see, e.g., Aurenhammer 1991 for an overview). Given a set of  $K$  points  $\{\hat{\xi}_1, \dots, \hat{\xi}_K\} \subset \Xi$ , the Voronoi diagram associated with these points defines a partition for each  $\hat{\xi}_i$  that contains only the  $\xi \in \Xi$  such that the Euclidean distance between  $\hat{\xi}_i$  and  $\xi$  is less than or equal to the distance to any other given point  $\hat{\xi}_j$ . We now apply this to the problem at hand: as  $\mathcal{A}$  is a finite set, we can express the partition induced by active uncertain parameter  $\hat{\xi}_i \in \mathcal{A}$  as

$$\begin{aligned} \Xi(\hat{\xi}_i) &= \Xi \cap \{\xi \mid \|\hat{\xi}_i - \xi\|_2 \leq \|\hat{\xi}_j - \xi\|_2 \quad \forall \hat{\xi}_j \in \mathcal{A}, \hat{\xi}_i \neq \hat{\xi}_j\} \\ &= \Xi \cap \left\{ \xi \mid (\hat{\xi}_j - \hat{\xi}_i) \cdot \xi \leq \frac{1}{2}(\hat{\xi}_j - \hat{\xi}_i) \cdot (\hat{\xi}_j + \hat{\xi}_i) \right. \\ &\quad \left. \forall \hat{\xi}_j \in \mathcal{A}, \hat{\xi}_i \neq \hat{\xi}_j \right\}, \end{aligned}$$

**Figure 1.** (Color online) Partitioning of simple uncertainty sets using Voronoi diagrams.

Notes. The first two diagrams use the polyhedral uncertainty set  $\Xi_p = \{\xi \mid \|\frac{1}{2}\xi_1, \xi_2\|_1 \leq 1\}$ , and the third uses an ellipsoidal set  $\Xi_e = \{\xi \mid \|\frac{1}{2}\xi_1, \xi_2\|_2 \leq 1\}$ . The bold points represent the active uncertain parameters  $\hat{\xi}$  used to create the partitions. In the center example, the red line represents an infinite set of possible active uncertain parameters, from which we have selected one arbitrarily.

which is a set defined by  $\Xi$  and  $|\mathcal{A}| - 1$  additional linear constraints (and if  $\Xi$  is polyhedral, then  $\Xi(\hat{\xi}_i)$  will also be polyhedral).

We will now generalize away from the cutting plane method-based reasoning and continuous problems, and we will assume throughout that the uncertainty set is closed and convex but not necessarily polyhedral. For constraint  $i$  and a static policy solution  $(\bar{x}^1, \bar{x}^2)$ , we define the set of active uncertain parameters for constraint  $i$  as

$$\mathcal{A}_i = \arg \min_{\xi \in \Xi} \{b_i(\xi) - \mathbf{a}_i^1(\xi) \cdot \bar{\mathbf{x}}^1 - \mathbf{a}_i^2(\xi) \cdot \bar{\mathbf{x}}^2\}, \quad (9)$$

which, as the set of optimal solutions for minimizing a linear function over a convex set is itself convex, may be a convex set of infinite cardinality. Note that we do not require that these  $\hat{\xi} \in \mathcal{A}_i$  correspond to constraints that have zero slack, as this may possibly never occur for a problem with integer variables. We can also define an equivalent set  $\mathcal{A}_c$  for the objective epigraph constraint. The possibility that these sets are of infinite cardinality presents a problem for our proposed Voronoi diagram scheme, as the natural generalization of Voronoi diagrams for convex sets (instead of just points) can produce nonconvex partitions (Lee and Drysdale 1981). We propose selecting a single point from each set  $\mathcal{A}_i$  using either problem-specific knowledge, or a more general technique such as taking the Chebyshev center or random selection. In Figure 1, we demonstrate the partitions that would be induced in a polyhedral uncertainty set for two different sets of uncertain parameters, as well as a partitioning of an ellipsoidal set.

## 2.2. Description of Method

We can now build an iterative method using this partitioning scheme. The method starts by solving a static policy version of our AO problem to determine a set of active uncertain parameters. We use these to construct a finitely adaptive version of our problem, and solve that. This, in turn, produces a new set of active uncertain parameters, which we can then use to partition further, ideally improving on the previous solution at each iteration.

We have already described how we collect a finite set of active uncertain parameters given a solution. The other key detail is the way we track the active uncertain parameters across iterations, and how we construct the partitions after the first iteration. We propose nested partitions: the

active uncertain parameters  $\mathcal{A}$  corresponding to the constraints for a partition at iteration  $k$  are used to subpartition that partition at iteration  $k + 1$ . We can describe the partition construction scheme in terms of a *tree*  $\mathcal{T}$  of uncertain parameters, for which we will define the following sets: *Leaves*( $\mathcal{T}$ ) is the set of leaves of the tree  $\mathcal{T}$ , *Children*( $\hat{\xi}$ ) is the set of children of  $\hat{\xi}$  in the tree  $\mathcal{T}$ , *Parent*( $\hat{\xi}$ ) is the parent of  $\hat{\xi}$  in the tree  $\mathcal{T}$ , and finally, *Siblings*( $\hat{\xi}$ ) is the set of children of the parent  $\hat{\xi}$ , i.e.,  $\text{Siblings}(\hat{\xi}) = \text{Children}(\text{Parent}(\hat{\xi}))$ . Each leaf of  $\mathcal{T}$  corresponds to a partition of the uncertainty set and each level of the tree corresponds to an iteration of the algorithm. At each iteration, we construct a partition for a leaf  $\hat{\xi}_i$  as an intersection of partitions

$$\begin{aligned} \Xi(\hat{\xi}_i) &= \{\xi \mid \|\hat{\xi}_i - \xi\|_2 \leq \|\hat{\xi}_j - \xi\|_2 \ \forall \hat{\xi}_j \in \text{Siblings}(\hat{\xi}_i)\} \\ &\cap \{\xi \mid \|\text{Parent}(\hat{\xi}_i) - \xi\|_2 \leq \|\hat{\xi}_j - \xi\|_2 \\ &\quad \forall \hat{\xi}_j \in \text{Siblings}(\text{Parent}(\hat{\xi}_i))\} \\ &\vdots \\ &\cap \Xi, \end{aligned} \quad (10)$$

which terminates when the parent is the root node, which has no siblings, and thus is equivalent to the entire uncertainty set. By adding the active uncertain parameters for partition  $\Xi(\hat{\xi}_i)$  as children of  $\hat{\xi}_i$  in  $\mathcal{T}$ , we create subpartitions of  $\Xi(\hat{\xi}_i)$  at the next iteration as shown in Figure 2.

Given this machinery, our method for solving problems of the form of (4) proceeds as follows.

### Partition-and-Bound Method for Two-Stage AMIO.

1. *Initialize.* Define  $\mathcal{T}^1$  to be the initial tree of uncertain parameters that consists of one root node (any  $\xi \in \Xi$ ). Set iteration counter  $k \leftarrow 1$ .

2. *Solve.* Solve the following partitioned version of (4), with one partition for every  $\hat{\xi}_j \in \text{Leaves}(\mathcal{T}^k)$ :

$$z_{\text{alg}}(\mathcal{T}^k) = \min_{\mathbf{x} \in \mathcal{X}, z}$$


subject to

$$\begin{aligned} \mathbf{c}^1(\xi) \cdot \mathbf{x}^1 + \mathbf{c}^2(\xi) \cdot \mathbf{x}_j^2 &\leq z \\ \forall \xi \in \Xi(\hat{\xi}_j), \ \forall \hat{\xi}_j \in \text{Leaves}(\mathcal{T}^k) \end{aligned} \quad (11)$$

$$\mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\xi) \cdot \mathbf{x}_j^2 \leq b_i(\xi) \quad \forall \xi \in \Xi(\hat{\xi}_j),$$

$$\forall \hat{\xi}_j \in \text{Leaves}(\mathcal{T}^k), \ i \in \mathcal{I},$$

**Figure 2.** (Color online) Visualization of the partitions and active uncertain parameter tree across multiple iterations.



*Note.* The uncertainty set is  $\Xi = \{(\xi_1, \xi_2) \mid |\xi_1 - 20|/10 + |\xi_2 - 10|/5 \leq 1\}$ , and we assume that at each iteration, all the extreme points of the set/partitions are active uncertain parameters (red points) in the solved partitioned problem.

where  $\Xi(\hat{\xi}_j)$  is the set defined in (10), and  $\mathbf{x}_j^2$  is the set of recourse decisions corresponding to the partition induced by parameters  $\hat{\xi}_j$ .

3. *Grow.* Initialize  $\mathcal{T}^{k+1} \leftarrow \mathcal{T}^k$ . For every leaf  $\hat{\xi}_j \in \text{Leaves}(\mathcal{T}^{k+1})$ , add children to that leaf for each  $\hat{\xi}$  in the set of active uncertain parameters  $\mathcal{A}$  for the partition  $\Xi(\hat{\xi}_j)$  and the solution to (11), selecting a finite subset if there is an infinite set of active uncertain parameters to choose from.

4. *Bound.* Calculate a lower bound  $z_{\text{lower}}(\mathcal{T}^{k+1})$  (see Section 2.3) of the fully adaptive solution, and terminate if the bound gap

$$\frac{(z_{\text{alg}} - z_{\text{lower}})}{|z_{\text{lower}}|} \quad (12)$$

is less than  $\epsilon_{\text{gap}}$ , or the iteration or time limit has been reached. Otherwise, set  $k \leftarrow k + 1$  and go to Step 2.

A possibility not directly addressed by this method is the case where the initial static solution is infeasible. For example, consider the feasibility problem

$$\begin{aligned} z_{\text{feas}} = \min_{x^2} & 0 \\ \text{subject to} & \quad x^2(\xi) \leq \xi + \epsilon \quad \forall \xi \in [0, 1] \\ & \quad x^2(\xi) \geq \xi - \epsilon \quad \forall \xi \in [0, 1], \end{aligned} \quad (13)$$

which has no feasible static policy (but does have the solution  $x^2(\xi) = \xi$ , as well as piecewise constant solutions). If infeasibility does occur, then one potential remedy is to

use problem-specific knowledge to identify a near-feasible solution or to initially relax problematic constraints so that active uncertain parameters can be collected for partitioning. Failing that, sampling (random or otherwise) of the uncertainty set can serve the same purpose. Once a feasible solution is identified, the algorithm can proceed as before.

Our method is similar to that of Postek and Den Hertog (2014) for the case of two-stage problems. The primary difference is the method of constructing partitions: while our method creates multiple child partitions at each iteration, the method of Postek and Den Hertog (2014) creates only two child partitions per iteration per partition. These two child partitions are defined by a single separating hyperplane between a pair of active uncertain parameters, with the primary selection heuristic employed in that paper being to choose the pair with the greatest distance between them. We will explore the implications of this difference further in Section 3 and the computational experiments.

We also note that, while we have focused here on mixed-integer linear optimization, other problem classes can be solved using the same approach. For example, we could consider quadratic terms in the objective, or second-order cone constraints. The only requirements we have for this method are the ability to extract active uncertain parameters after solving a partitioned problem to further improve the solution at the next iteration, and the ability to partition the uncertainty set.

### 2.3. Calculating Lower Bounds

Our method ideally produces successfully closer approximations to the fully adaptive problem (4). The most natural termination criterion then, apart from an iteration or time limit criterion, is to terminate when our approximation (an upper bound, as we are minimizing) is sufficiently close to the fully adaptive solution. As we do not know the fully adaptive solution, we propose using the set of uncertain parameters available at the end of iteration  $k$  to obtain a lower bound. That is, after the solution of the partitioned problem at iteration  $k$ , we grow our tree  $\mathcal{T}$  and use all parameters in  $\mathcal{T}$  to construct and solve a deterministic problem whose objective provides the lower bound. This is essentially the same as the “scenario-based bound” of Hadjiyiannis et al. (2011), which also uses “binding” uncertainty realizations, much like our active uncertain parameters. We restate this bound here in the context of our method.

**PROPOSITION 1.** *Consider the solution of the deterministic optimization problem*

$$\begin{aligned} z_{\text{lower}}(\mathcal{T}) &= \min_{\mathbf{x} \in \mathcal{X}, z} \\ \text{subject to} \quad & \mathbf{c}^1(\hat{\xi}_j) \cdot \mathbf{x}^1 + \mathbf{c}^2(\hat{\xi}_j) \cdot \mathbf{x}_j^2 \leq z \quad \forall \hat{\xi}_j \in \mathcal{T} \\ & \mathbf{a}_i^1(\hat{\xi}_j) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\hat{\xi}_j) \cdot \mathbf{x}_j^2 \leq b_i(\hat{\xi}_j) \quad \forall \hat{\xi}_j \in \mathcal{T}, i \in I, \end{aligned} \quad (14)$$

where  $\mathcal{T}$  is a set of uncertain parameters  $\hat{\xi} \in \Xi$ . Then,  $z_{\text{lower}}(\mathcal{T}) \leq z_{\text{full}}$ , where  $z_{\text{full}}$  is the solution to (4).

**PROOF OF PROPOSITION 1.** Problem (14) is a relaxation of (4) as we only require the solution to satisfy the constraints for a finite subset of uncertain parameters. Any solution to the fully adaptive problem (4) has a trivial mapping to a feasible solution for (14) that preserves the objective function value, so  $z_{\text{lower}}(\mathcal{T}) \leq z_{\text{full}}$ .  $\square$

A useful property of this bound is that as the tree grows the lower bound will tend to approach the fully adaptive solution. At the same time, we are improving our upper bound by decreasing the size of each partition, and thus closing the bound gap from both sides. Care must be taken if solving (14) approximately, for example, if (14) is a MIO problem and we allow the solver to terminate when the gap between the best-known integer solution and the continuous relaxation falls below a tolerance, as is common in practice, we will obtain a solution that is larger than the true lower bound or even the fully adaptive solution, and thus understate the bound gap.

### 2.4. Worked Example: Two-Stage Inventory Control Problem

**EXAMPLE 1.** We now demonstrate our approach by solving a two-stage inventory control problem. We must satisfy all of the uncertain demand, and there are three ways to order

stock to do so: ordering any amount here and now at a unit cost of 50, ordering a fixed “lot” of size 25 at a unit cost of 60, and ordering a fixed lot of size 25 at a unit cost of 75. Only one lot of each size may be ordered, but the decision to order them can be made after the demand is known. We must pay a high unit cost of 65 for the storage and disposal of any of the remaining stock after the demand is satisfied. We can formulate this as a two-stage AMIO problem

$$\begin{aligned} \min \quad & z \\ \text{subject to} \quad & 50x^1 + 65I^2(\xi) + 1,500y_A^2(\xi) \\ & + 1,875y_B^2(\xi) \leq z \quad \forall \xi \in \Xi \\ & I^2(\xi) \geq 0 \quad \forall \xi \in \Xi \\ & x^1 \geq 0, y_A^2(\xi), y_B^2(\xi) \in \{0, 1\} \quad \forall \xi \in \Xi, \end{aligned} \quad (15)$$

where  $x^1$  is the here-and-now-continuous ordering decision,  $y_A^2(\xi)$  and  $y_B^2(\xi)$  are the wait-and-see binary ordering decisions,  $I^2(\xi) = x^1 - \xi + 25y_A^2(\xi) + 25y_B^2(\xi)$ , and  $\xi$  is the uncertain demand that is drawn from the uncertainty set  $\Xi = \{\xi \mid 5 \leq \xi \leq 95\}$ . We will initialize our tree  $\mathcal{T}^1$  (Step 1) with a root  $\hat{\xi} = 50$ , leading to a problem with a single partition (the entire set) and with  $y_A^2$  and  $y_B^2$  unable to vary with  $\xi$  (static). Solving this single partition version of (15) (Step 2) gives us the initial solution  $z = 10,600$ ,  $x^1 = 95$ ,  $y_A^2 = 0$ , and  $y_B^2 = 0$ . The active uncertain parameter value for the objective constraint is when the demand is low ( $\hat{\xi} = 5$ , which leaves 90 units of stock) and the worst case for the nonnegative inventory constraint is when the demand is high ( $\hat{\xi} = 95$ , leaving no stock). We grow the tree (Step 3) to obtain  $\mathcal{T} = \{50: \{5, 95\}\}$ , which we then use to determine a lower bound on the fully adaptive solution ( $z_{\text{lower}} = 5,625$ ) for an initial gap of approximately 88%.

We now solve the two partition version of (15) at iteration 2. The first partition is

$$\begin{aligned} \Xi(\hat{\xi}_1 = 5) &= \{\xi \mid 5 \leq \xi \leq 95\} \cap \{\xi \mid \|5 - \xi\|_2 \leq \|95 - \xi\|_2\} \\ &= \{\xi \mid 5 \leq \xi \leq 50\}, \end{aligned}$$

and the second is  $\Xi(\hat{\xi}_2 = 95) = \{\xi \mid 50 \leq \xi \leq 95\}$ . The objective function value for this problem is substantially lower ( $z = 7,925$ ) than at iteration 1 and the stock bought initially is lower ( $x^1 = 70$ ). For the first partition, we will order neither lot ( $y_{A,1}^2 = y_{B,1}^2 = 0$ ), but if the demand falls in the second partition, we will purchase the first, cheaper lot ( $y_{A,2}^2 = 1, y_{B,2}^2 = 0$ ). Once again, the active uncertain parameters lie at the extremes of the partitions, which leads to the tree  $\mathcal{T}^3 = \{50: \{5: \{5, 50\}, 95: \{50, 95\}\}\}$ . Our lower bound is unchanged, as we were already using the values 5, 50, and 95, but our gap shrinks substantially to 41%, thanks to the reduced upper bound.

We now begin our third and final iteration, by solving the partitioned problem induced by  $\mathcal{T}^3$ . In the interest of space, we only detail the construction of the partition for the leaf  $\hat{\xi} = 50$  that is the child of  $\hat{\xi} = 5$ :

$$\begin{aligned} \Xi(\hat{\xi} = 50) &= \Xi \cap \{\xi \mid \|5 - \xi\|_2 \leq \|95 - \xi\|_2\} \\ &\cap \{\xi \mid \|50 - \xi\|_2 \leq \|5 - \xi\|_2\} \end{aligned}$$



$$\begin{aligned} &= \Xi \cap \{\xi \mid \xi \leq 50\} \cap \{\xi \mid \xi \geq 27\frac{1}{2}\} \\ &= \{\xi \mid 27\frac{1}{2} \leq \xi \leq 50\}. \end{aligned}$$

The other three partitions are  $\{\xi \mid 5 \leq \xi \leq 27\frac{1}{2}\}$ ,  $\{\xi \mid 50 \leq \xi \leq 72\frac{1}{2}\}$ , and  $\{\xi \mid 72\frac{1}{2} \leq \xi \leq 95\}$ . Solving this four-partition problem produces the best solution yet ( $z = 7,375$ ). The solution purchases substantially less stock up front ( $x^1 = 47\frac{1}{2}$ ), and we can best describe the values of  $y_A$  and  $y_B$  as piecewise constant functions of  $\xi$ , i.e.,

$$y_A^2(\xi) = \begin{cases} 0, & 5 \leq \xi < 27\frac{1}{2}, \\ 1, & 27\frac{1}{2} \leq \xi \leq 95, \end{cases} \quad (16)$$

and

$$y_B^2(\xi) = \begin{cases} 0, & 5 \leq \xi < 27\frac{1}{2}, \\ 1, & 27\frac{1}{2} \leq \xi \leq 95, \end{cases} \quad (17)$$

respectively. Growing the tree with the new active uncertain parameters leads to a new lower bound of  $z_{\text{lower}} = 5,912\frac{1}{2}$ , resulting in a final gap of 26%. Although it cannot be determined at this iteration, our upper bound is very good as the fully adaptive solution has an objective of  $z_{\text{full}} = 7,250$  (as determined by exhaustively applying this method).

## 2.5. Incorporating Affine Adaptability

In Section 1, we discussed the popularity and success of affine adaptability, also known as linear decisions rules. We can express this type of adaptability in the context of our two-stage problem (4) with the substitution

$$\mathbf{x}^2(\xi) = \mathbf{F}\xi + \mathbf{g}, \quad (18)$$

where  $\mathbf{F}$  and  $\mathbf{g}$  are now the decision variables instead of  $\mathbf{x}^2$  itself. There are two problems with substituting (18) directly into (4). The first is that if  $\mathbf{a}_i^2(\xi)$  and  $\mathbf{c}^2(\xi)$  are not constant with respect to the uncertain parameters, then we will have quadratic products of uncertain parameters, which is not generally tractable (Ben-Tal et al. 2004). The second problem is that if a variable  $x_j^2(\xi)$  is integer, then  $\mathbf{F}_j$  must necessarily be zero for any nontrivial uncertainty set and so there is no adaptability.

We will thus separate our second-stage decisions into two sets: continuous decisions  $\mathbf{x}^2$  and integer decisions  $\mathbf{y}^2$ , and apply our affine substitution only to the continuous decisions. This leads to an affine adaptability form of the fully adaptive problem (4):

$$\begin{aligned} &\min_{\mathbf{F}, \mathbf{g}, \mathbf{x}^1, \mathbf{y}, z} z \\ &\text{subject to } \mathbf{c}^1(\xi) \cdot \mathbf{x}^1 + \mathbf{c}_x^2 \cdot (\mathbf{F}\xi + \mathbf{g}) \\ &\quad + \mathbf{c}_y^2(\xi) \cdot \mathbf{y}^2(\xi) \leq z \quad \forall \xi \in \Xi \\ &\quad \mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_{x,i}^2 \cdot (\mathbf{F}\xi + \mathbf{g}) + \mathbf{a}_{y,i}^2(\xi) \cdot \mathbf{y}^2(\xi) \\ &\quad \leq b_i(\xi) \quad \forall \xi \in \Xi, i \in \{1, \dots, m\} \\ &\quad \mathbf{F} \in \mathbb{R}^{n_x \times n_\xi}, \mathbf{g} \in \mathbb{R}^{n_x}, \mathbf{y} \in \mathbb{Z}^{n_y}. \end{aligned} \quad (19)$$

A key observation with respect to our finite partitioning method is that  $\mathbf{F}$  and  $\mathbf{g}$  are themselves decisions that can

be deferred until the uncertainty has been realized, and so they too are second-stage variables. This allows us to combine finite adaptability and affine adaptability, as we can associate a different affine policy  $(\mathbf{F}_j, \mathbf{g}_j)$  for each partition. Doing so produces a piecewise linear policy, for example, for two partitions, we would have

$$\mathbf{x}^2(\xi) = \begin{cases} \mathbf{F}_1\xi + \mathbf{g}_1, & \xi \in \Xi(\hat{\xi}_1), \\ \mathbf{F}_2\xi + \mathbf{g}_2, & \xi \in \Xi(\hat{\xi}_2), \end{cases} \quad (20)$$

for continuous second-stage decisions, and a piecewise constant policy

$$\mathbf{y}^2(\xi) = \begin{cases} \mathbf{y}_1^2, & \xi \in \Xi(\hat{\xi}_1), \\ \mathbf{y}_2^2, & \xi \in \Xi(\hat{\xi}_2), \end{cases} \quad (21)$$

for integer decisions. We demonstrate combining affine adaptability with finite adaptability in our multistage lot-sizing computational example in Section 5.2, including a comparison against the method of Postek and Den Hertog (2014), which can be used to create a similar combination.

## 3. Analysis of the Partition-and-Bound Method

In this section, we characterize the partition-and-bound method proposed above from a theoretical point of view. In particular, we seek to understand how the algorithm approaches the ideal fully adaptive solution, and to what degree the number of partitions grows at each iteration.

### 3.1. Convergence Properties

We now aim to understand how  $z_{\text{alg}}(\mathcal{T}^k)$  varies with the iteration count  $k$ .

**PROPOSITION 2.** *The upper bound  $z_{\text{alg}}(\mathcal{T}^k)$  will never increase as  $k$  increases.*

**PROOF OF PROPOSITION 2.** This follows from the “nested” nature of our partitions. Consider the solution of the partitioned problem at iteration  $k$ , e.g.,  $\bar{\mathbf{x}}^1, \bar{\mathbf{x}}_1^2, \bar{\mathbf{x}}_2^2, \dots$ . At iteration  $k + 1$ , we will further subpartition each of these partitions, resulting in new second-stage variables. We can construct a feasible solution for iteration  $k + 1$  by setting the value of each of these new variables for each subpartition to the same value as the decision variables associated with their respective parent partition. This new solution is feasible by construction and has the same objective value, so we will at least match the objective of iteration  $k$  at iteration  $k + 1$ .  $\square$

This property has computational implications, as it allows us to provide an initial feasible solution at iteration  $k + 1$  using the solution at iteration  $k$ . This can allow integer optimization solvers to truncate portions of the search tree faster, and reach lower integrality gaps earlier.

We have established that we will never get worse, and it is also possible to identify families of problems (in particular, problems with only continuous decision variables), where we will smoothly converge to the fully adaptive solution. Instead of identifying these, we instead provide the following counterexample that demonstrates that there does not exist a general guarantee that we will improve the bound gap for any finite number of iterations, even if we start at far from the lower bound, for linear AMIO problems. To demonstrate this, consider the parametric family of problems

$$\begin{aligned} z(\epsilon) = & \min_{x^2 \in \{0, 1\}, y^2 \in \{0, 1\}, z} z \\ \text{subject to } & x^2(\xi) + y^2(\xi) \leq z \quad \forall \xi \in [0, 1] \\ & x^2(\xi) \geq \frac{\epsilon - \xi}{\epsilon} \quad \forall \xi \in [0, 1] \\ & y^2(\xi) \geq \frac{-\epsilon + \xi}{\epsilon} \quad \forall \xi \in [0, 1], \end{aligned} \quad (22)$$

where  $\epsilon \in [0, 1]$ . This problem has a fully adaptive solution of  $z = 1$  and

$$x^2(\xi) = \begin{cases} 1, & 0 \leq \xi \leq \epsilon, \\ 0, & \epsilon < \xi \leq 1, \end{cases} \quad y^2(\xi) = \begin{cases} 0, & 0 \leq \xi \leq \epsilon, \\ 1, & \epsilon < \xi \leq 1. \end{cases}$$

Note that the static policy sets  $x^2 = y^2 = 1$ , and thus  $z = 2$ . The initial bound gap is 1 as any random  $\hat{\xi}$  used to calculate the lower bound will correctly identify it to be  $z_{\text{lower}} = 1$ . Both variables share the same “breakpoint” of  $\epsilon$ , which means a simple two partition finitely adaptive solution could match the fully adaptive solution, if only such a partition could be identified.

**PROPOSITION 3.** *The partition-and-bound method will never find a solution equivalent to the fully adaptive solution for (22), or decrease the bound gap, unless  $\epsilon = q \times 2^{-p}$  for some  $p, q \in \mathbb{Z}^+$ .*

**PROOF OF PROPOSITION 3.** First, note that the two extreme points of each partition are always the active uncertain parameters (the low end from the second constraint, and the high end from the third). Thus, at the beginning of iteration  $k$ , we will have  $2^{k-1}$  partitions of width  $2^{1-k}$ . For all partitions  $\Xi(\xi) = [a, b]$  such that  $b < \epsilon$  the corresponding decision variable values are  $x^2 = 1$  and  $y^2 = 0$ , with the corresponding lower bound of 1 on the objective  $z$ . For all partitions such  $a > \epsilon$ , we have decision variable values  $x^2 = 0$  and  $y^2 = 1$ , also with the corresponding lower bound of 1 on the objective  $z$ . The remaining partition then has the property that  $a \leq \epsilon \leq b$ . If the inequalities are strict, then we must set  $x^2 = y^2 = 1$ , and thus do not improve over the static policy. However, if  $\epsilon = a$  or  $\epsilon = b$ , then we can match the fully adaptive solution because the “breakpoint” is correct. As  $a$  and  $b$  are always of the form  $q \times 2^{-p}$  for nonnegative integers  $p$  and  $q$ , this condition can only be satisfied if  $\epsilon$  is also drawn from this family. Any other

choice of  $\epsilon$ , such as  $1/3$ , will never be matched for any finite number of iterations.  $\square$

This result also applies to the method of Postek and Den Hertog (2014), using similar reasoning. While this result is a negative one, we believe that such structures are uncommon in problems of interest, and may be avoidable by making different modeling choices when they do occur.

### 3.2. Growth in Partitioning Scheme

For a problem with  $m$  uncertain constraints and an uncertain objective, the number of partitions with our method is at most  $(m+1)^{k-1}$  at iteration  $k$  (and may be much lower if the same uncertain parameter is active for multiple constraints). In contrast, the proposal of Postek and Den Hertog (2014) will have at most  $2^{k-1}$ , but may require more iterations to achieve a solution of the same quality. Here, we show by example that the growth rate of the size of the partitioning scheme is of secondary concern to the relationship between computational effort and quality of solution for many cases.

Consider the following example problem, which was presented in Bertsimas and Goyal (2010):

$$\begin{aligned} z(n) = & \min_{x^2 \geq 0, z} z \\ \text{subject to } & \mathbf{e} \cdot \mathbf{x}^2(\xi) \leq z \quad \forall \xi \in \Xi \\ & \mathbf{x}^2(\xi) \geq \xi \quad \forall \xi \in \Xi, \end{aligned} \quad (23)$$

where  $\Xi = \{\xi \mid \sum_{i=1}^n \xi_i \leq 1, \xi \geq 0\}$ . The initial static policy has an objective value of  $n$  compared to the fully adaptive policy objective value of 1. The upper bound provided by our method in the second iteration is  $(n+1)/2$ , which is obtained using  $n$  partitions. On the other hand, the method of Postek and Den Hertog (2014) will have  $2^{k-1}$  partitions at iteration  $k$ . If we assume that the time to solve this problem is proportional to the number of partitions, then in the same time it takes to solve the second iteration for our method ( $n$  time units), we can do approximately  $\log_2(n)$  iterations of the Postek and Den Hertog (2014) method. However, we can show (Appendix A) that it requires  $n$  iterations of that method (and thus  $2^{n-1}$  partitions) to achieve the same objective as our method, demonstrating the need to focus on the solution quality versus time rather than iteration or partition counts.

In the above example, all partitions had the same objective contribution, i.e., the overall objective was equal to the objective value for each partition. This is not necessarily always the case, and it is trivial to find cases where this does not occur. Recall that the motivation for our method, as presented in Theorem 1, was to identify active uncertain parameters that were “binding” the solution and to separate them. A similar approach can be taken with partitions themselves: if a partition’s objective is not restricting the overall objective (the partition is not *active*), then we will not subpartition it at the next iteration. To identify these

partitions, we need to encourage a solver to find the best-possible solution for each partition. We can achieve this by adding a very small term to the objective, e.g.,

$$z_{\text{alg}}(\mathcal{T}^k) = \min_{\mathbf{x} \in \mathcal{X}, z} z + \epsilon \sum_j \tilde{z}_j$$

subject to

$$\begin{aligned} \tilde{z}_j &\leq z \quad \forall \hat{\xi}_j \in \text{Leaves}(\mathcal{T}^k) \\ \mathbf{c}^1(\xi) \cdot \mathbf{x}^1 + \mathbf{c}^2(\xi) \cdot \mathbf{x}_j^2 &\leq \tilde{z}_j \\ &\quad \forall \xi \in \Xi(\hat{\xi}_j), \forall \hat{\xi}_j \in \text{Leaves}(\mathcal{T}^k) \\ \mathbf{a}_i^1(\xi) \cdot \mathbf{x}^1 + \mathbf{a}_i^2(\xi) \cdot \mathbf{x}_j^2 &\leq b_i(\xi) \\ &\quad \forall \xi \in \Xi(\hat{\xi}_j), \forall \hat{\xi}_j \in \text{Leaves}(\mathcal{T}^k), i \in \mathcal{I}, \end{aligned} \quad (24)$$

where  $\tilde{z}_j$  is the objective for partition  $j$ , and we say a partition is *active* if  $\tilde{z}_j = z$ .

Consider the following variation on (24) that weights one of the components of the decision above the rest:

$$\begin{aligned} z(M, n) &= \min_{\mathbf{x}^2 \geq 0, z} z \\ \text{subject to} \quad (M-1)x_1^2 + \mathbf{e} \cdot \mathbf{x}^2(\xi) &\leq z \\ &\quad \forall \xi \in \Xi \\ \mathbf{x}^2(\xi) &\geq \xi \quad \forall \xi \in \Xi, \end{aligned} \quad (25)$$

where the uncertainty set is the same as in the previous example. The static policy now has an objective value of  $M + n - 1$  compared to the fully adaptive solution's  $M$ . If we now apply our partitioning scheme once, the partition corresponding to  $i = 1$  will have objective  $M + (n - 1)/2$  and all other partitions will have the smaller objective  $(M + n)/2$ . At the next iteration, this is further reduced to  $M + (n - 1)/4$  for the active partition and  $(M + n + 2)/4$  for all other partitions. Consider the specific case of  $M = n = 10$ : the initial static solution has objective 19. After partitioning, we have one partition with objective 14.5 and the remainder have objective 10, which is equal to the fully adaptive objective. Thus we should only add children parameters for the active partition, as further subpartitions of the inactive partitions will not have any further effect on the overall objective. If we continue with further iterations, we find that the number of partitions grows linearly with the number of iterations in this case. We implement this approach in our computation experiments in Section 5.

#### 4. Partition-and-Bound Method for Multistage AMIO

We now generalize our partitioning method described above for two-stage AMIO to the full multistage case of (2). In the interests of clarity, the order of events in our model is as follows: The here-and-now decision  $\mathbf{x}^1$  are made before any uncertainty is realized, and then the first-stage uncertain parameters  $\xi^1$  are revealed. We make the decision  $\mathbf{x}^2$  with the benefit of knowing  $\xi^1$ , but with no other knowledge of the uncertainty parameters to be

revealed later except what is implied by knowing  $\xi^1$ . This then proceeds until we make the decision  $\mathbf{x}^T$  with the benefit of knowing  $\xi^1, \dots, \xi^{T-1}$ , and the final uncertain parameters  $\xi^T$  are then revealed (if applicable). Dependencies between uncertain parameters within and across time stages can be modeled in the uncertainty set with constraints.

While this model of decision making across time is simple enough to state, modeling these *nonanticipativity restrictions*—that a decision made now cannot be made by using exact knowledge of the future—is the primary complication that we address in this section as we extend our approach to the multistage case. We note that enforcing these restrictions is trivial for some classes of adaptability. For example, for affine adaptive policies, we can restrict the policy to be a function of only the appropriate components of  $\xi$ . It is less obvious how to proceed with finite adaptability. The following example demonstrates how a naive application of the method developed in Section 2 will either produce overly optimistic solutions if we fail to recognize a need for nonanticipativity constraints, or will produce pessimistic solutions where the choice of partitioning scheme and nonanticipativity constraints combine to result in less adaptability than we may like.

**EXAMPLE 2.** Consider a three-stage ( $T = 3$ ) version of the inventory problem of Section 2.4. The problem proceeds exactly as before, except there are two rounds of demand, two rounds of continuous order decisions, and two rounds of lot ordering decisions. The holding cost is paid on the stock remaining after rounds of demand and all parameters are the same as before, resulting in the three-stage AMIO


$$\begin{aligned} \min \quad & z \\ \text{subject to} \quad & 50x^1 + 65I^2(\xi^1) + 1,500y_A^2(\xi^1) \\ & \quad + 1,875y_B^2(\xi^1) + 50x^2(\xi^1) \\ & \quad + 65I^3(\xi^1, \xi^2) + 1,500y_A^3(\xi^1, \xi^2) \\ & \quad + 1,875y_B^3(\xi^1, \xi^2) \leq z \quad \forall \xi \in \Xi \\ & I^2(\xi^1), I^3(\xi^1, \xi^2), x^1, x^2(\xi^1) \geq 0 \\ & \quad \forall \xi \in \Xi \\ & y_A^2(\xi^1), y_B^2(\xi^1), y_A^3(\xi^1, \xi^2), \\ & \quad y_B^3(\xi^1, \xi^2) \in \{0, 1\} \quad \forall \xi \in \Xi, \end{aligned} \quad (26)$$

where

$$\begin{aligned} I^2(\xi^1) &= x^1 - \xi^1 + 25y_A^2(\xi^1) + 25y_B^2(\xi^1), \\ I^3(\xi^1, \xi^2) &= I^2(\xi^1) + x^2(\xi^1) - \xi^2 \\ & \quad + 25y_A^3(\xi^1, \xi^2) + 25y_B^3(\xi^1, \xi^2), \end{aligned}$$

and  $\Xi = \{5 \leq \xi^1, \xi^2 \leq 95\}$ . We begin by solving the static policy version of this problem (no partitions), and obtain a solution with objective value 27,050 and three active uncertain parameters:  $\hat{\xi}_1 = (5, 5)$  (for the objective constraint),  $\hat{\xi}_2 = (95, 5)$  (for the nonnegativity constraint on  $I^2$ ), and  $\hat{\xi}_3 = (95, 95)$  (for the nonnegativity constraint on  $I^3$ ). We now construct partitions exactly as before, ignoring the

**Figure 3.** (Color online) Partitioning of the uncertainty set in Example 2 using the two-stage partitioning scheme without any adjustments for the multistage case.



*Note.* The thin solid lines for  $\xi^1 = 15$  and  $\xi^1 = 85$  demonstrate that knowing components  $\xi^1$  and  $\xi^2$  is required to fully identify a partition, requiring the addition of nonanticipativity constraints that remove many of the benefits of partitioning scheme.

multistage nature of the uncertain parameters, resulting in the three partitions:

$$\begin{aligned}\Xi(\hat{\xi}_1 = (5, 5)) &= \{\xi \mid 5 \leq \xi^1 \leq 50, \xi^1 + \xi^2 \leq 100\}, \\ \Xi(\hat{\xi}_2 = (95, 5)) &= \{\xi \mid 50 \leq \xi^1 \leq 95, 5 \leq \xi^2 \leq 50\}, \\ \Xi(\hat{\xi}_3 = (95, 95)) &= \{\xi \mid 5 \leq \xi^1 \leq 95, \xi^1 + \xi^2 \geq 100\},\end{aligned}$$

which are visualized in Figure 3. As this is a multistage problem we now consider nonanticipativity. We proceed by creating a set of variables for each partition, for example,  $x_1^2$ ,  $x_2^2$ , and  $x_3^2$  for the second-stage decision  $x^2(\xi^1)$ . Recall that the interpretation of finite adaptability is that if the realized value of  $\xi$  falls in  $\Xi(\hat{\xi}_1)$ , then we will select  $x_1^2$  as the implemented decision, and similarly for the other partitions. With this in mind, we consider the case when  $\xi^1 = 15$ . This value of  $\xi^1$  may fall in either  $\Xi(\hat{\xi}_1)$  or  $\Xi(\hat{\xi}_3)$ , but we cannot determine which decision to implement without knowing  $\xi^2$ , which is yet to occur. We must then add the nonanticipativity restriction that  $x_1^2 = x_3^2$ . Similar logic applies to  $\xi^1 = 85$ , which may be in either  $\Xi(\hat{\xi}_2)$  or  $\Xi(\hat{\xi}_3)$ , which requires us to add the restriction  $x_2^2 = x_3^2$ . These restrictions combine to leave us with no adaptability at all in our second-stage decisions ( $x_1^2 = x_2^2 = x_3^2$ ). The third-stage variables are not affected by these problems as, by that time, we know  $\xi^1$  and  $\xi^2$ , and can thus unambiguously identify the partition and the decision we will implement.

#### 4.1. Multistage Partitions and Nonanticipativity Constraints

A naive application of our existing partitioning scheme will result in overly conservative solutions after the application of nonanticipativity constraints. We thus seek a minimal modification of our existing scheme that retains as much of the

benefit of partitioning as possible when these constraints are enforced. We will again have a partition for each leaf of our tree  $\mathcal{T}$ , and we will use the same tree operators and terminology we established previously. The changes are in how we convert the leaves into partitions, and are coupled with how we identify the required nonanticipativity constraints.

Our scheme is very similar to the two-stage version presented previously, but with restrictions at each layer of the tree to only certain elements of the uncertain parameters. In particular, the partition induced by a leaf  $\hat{\xi}_i$  is defined to be

$$\begin{aligned}\Xi(\hat{\xi}_i) &= \{\xi \mid \|\hat{\xi}_i^{t_{i,j}} - \xi^{t_{i,j}}\|_2 \leq \|\hat{\xi}_j^{t_{i,j}} - \xi^{t_{i,j}}\|_2 \\ &\quad \forall \hat{\xi}_j \in \text{Siblings}(\hat{\xi}_i)\} \\ &\cap \{\xi \mid \|\text{Parent}(\hat{\xi}_i)^{t'_{i,j}} - \xi^{t'_{i,j}}\|_2 \leq \|\hat{\xi}_j^{t'_{i,j}} - \xi^{t'_{i,j}}\|_2 \\ &\quad \forall \hat{\xi}_j \in \text{Siblings}(\text{Parent}(\hat{\xi}_i))\} \\ &\cap \dots \cap \Xi.\end{aligned}\quad (27)$$

We define  $t_{i,j}$  for a pair of uncertain parameters  $\hat{\xi}_i$  and  $\hat{\xi}_j$  as  $\arg \min_t \{\hat{\xi}_i^t \neq \hat{\xi}_j^t\}$ , that is, the first time stage for which the uncertain parameters differ.  $t'_{i,j}$  is the same but for the one level up the tree, i.e., the first time stage  $t$  such that  $\text{Parent}(\hat{\xi}_i)^t \neq \text{Parent}(\hat{\xi}_j)^t$  for  $\hat{\xi}_i, \hat{\xi}_j \in \text{Siblings}(\text{Parent}(\hat{\xi}_i))$ , and so on.

We must also present a complementary scheme for enforcing nonanticipativity. We claim the following for two partitions of the uncertainty set  $\Xi(\hat{\xi}_i)$  and  $\Xi(\hat{\xi}_j)$  induced by the leaf uncertain parameters  $\hat{\xi}_i$  and  $\hat{\xi}_j$ .

**PROPOSITION 4.** *If there exists  $\psi = (\psi^1, \dots, \psi^T) \in \Xi(\hat{\xi}_i)$  and  $\phi = (\phi^1, \dots, \phi^T) \in \Xi(\hat{\xi}_j)$  such that  $\psi^s = \phi^s \forall s \in \{1, \dots, t-1\}$ , and at least one of  $\psi \in \text{int}(\Xi(\hat{\xi}_i))$  and  $\phi \in \text{int}(\Xi(\hat{\xi}_j))$  holds, then we must enforce nonanticipativity constraints for the corresponding decisions at time stage  $t$  as the two partitions cannot be distinguished with the uncertain parameters realized by that time stage. Otherwise, we do not need to enforce any restrictions at time stage  $t$  for this pair.*

**PROOF OF PROPOSITION 4.** If the condition holds, then there exists a partial realization of the uncertain parameters, which could lie in two different partitions, based on known information, so we must restrict the decisions for those partitions to be the same. If the equality condition holds, but only for a point on the boundary of both partitions, then both decisions are valid and interchangeable, and we gain no benefit from future knowledge separating the two anyway. If the equality condition does not hold, then the partitions are distinguishable and no constraint is required.  $\square$

Given this condition, we can now propose a simple routine for determining which constraints to add. For the sake of exposition, we will focus on polyhedral uncertainty sets, although a similar routine could be constructed for more general sets. In particular, for each pair of leaf parameters  $\hat{\xi}_i$  and  $\hat{\xi}_j$  and their corresponding induced partitions, expressed as

$$\Xi(\hat{\xi}_i) = \{\xi \in \Xi \mid \mathbf{F}\xi \leq \mathbf{f}\} \quad \text{and} \quad \Xi(\hat{\xi}_j) = \{\xi \in \Xi \mid \mathbf{G}\xi \leq \mathbf{g}\},$$



we can formulate the optimization problem

$$\begin{aligned} z(t) = & \max_{\psi, \phi, \alpha \geq 0, \beta \geq 0} \{\alpha + \beta\} \\ \text{subject to } & F\psi + \alpha \mathbf{e} \leq \mathbf{f} \\ & G\phi + \beta \mathbf{e} \leq \mathbf{g} \\ & \psi^s = \phi^s \quad \forall s \in \{1, \dots, t-1\}, \end{aligned} \quad (28)$$

the solution to which informs us if nonanticipativity constraints are necessary at time stage  $t$ . If the problem is infeasible, then they are not required, as the first two of constraints can always be satisfied, but the third may not necessarily be satisfiable. If there is a feasible solution but the objective value is 0, then all feasible solutions are on the boundaries of the uncertainty sets, which does not require the constraints. If the objective value is positive, then we must enforce the nonanticipativity constraints as there is an interior point for the partitions that is identical up to time  $t$ .

Finally, we present a simple alternative rule by which the need for nonanticipativity constraints can be determined. This rule is not a general method for enforcing nonanticipativity in finite adaptability, as it relies on the particular structure produced by our partitioning scheme. Given two partitions,  $\Xi(\hat{\xi}_i)$  and  $\Xi(\hat{\xi}_j)$ , the rule for whether we need to enforce the nonanticipativity constraints  $\mathbf{x}_i^t = \mathbf{x}_j^t$  is as follows:

1. Let  $\hat{\xi}_A \leftarrow \hat{\xi}_i$  and  $\hat{\xi}_B \leftarrow \hat{\xi}_j$ .
2. If  $\hat{\xi}_A$  and  $\hat{\xi}_B$  are siblings, then let  $t^{A,B} = \arg \min_s \{\hat{\xi}_A^s \neq \hat{\xi}_B^s\}$ . If  $t > t^{A,B}$ , then we *do not* need to enforce the nonanticipativity constraints. If  $t \leq t^{A,B}$ , then we *do*.
3. Otherwise, if  $\hat{\xi}_A$  and  $\hat{\xi}_B$  are not siblings, let  $\hat{\xi}_A \leftarrow \text{Parent}(\hat{\xi}_A)$  and  $\hat{\xi}_B \leftarrow \text{Parent}(\hat{\xi}_B)$ , and go to Step 2.

This rule is simple to implement, but can be overconservative as it may declare nonanticipativity constraints as required when they are not needed. Despite this shortcoming, it could be used to quickly determine the cases where the constraints are not required before the above optimization approach can definitively determine whether

they are truly required. We provide proof of the correctness of the rule, as well as an example demonstrating overconservatism, in Appendix B. This scheme, applied to Example 2, is shown in Figure 4.

## 4.2. Multistage Method and Implementation


In the interests of space, we will not present the multistage method in full, as it suffices to simply take the definition of the method for the two-stage case, but use the multistage partitioning scheme and add nonanticipativity constraints wherever they are required. We can also extend our lower-bound approach to the multistage case, in particular,

$$\begin{aligned} z_{\text{lower}}(\mathcal{T}) = & \min_{\mathbf{x} \in \mathcal{X}, z} z \\ \text{subject to } & \sum_{t=1}^T \mathbf{c}'(\hat{\xi}_j) \cdot \mathbf{x}_j^t \leq z \quad \forall \hat{\xi}_j \in \mathcal{T} \\ & \sum_{t=1}^T \mathbf{a}'_i(\hat{\xi}_j) \cdot \mathbf{x}_j^t \leq b_i(\hat{\xi}_j) \\ & \forall \hat{\xi}_j \in \mathcal{T}, i \in I \\ & \mathbf{x}_j^t = \mathbf{x}_k^t \quad \forall \hat{\xi}_j, \hat{\xi}_k \in \mathcal{T}, \\ & \forall t: \hat{\xi}_j^{1,\dots,t-1} = \hat{\xi}_k^{1,\dots,t-1}, \end{aligned} \quad (29)$$

where the final set of constraints enforces the nonanticipativity restrictions. If two samples  $\hat{\xi}_j$  and  $\hat{\xi}_k$  are identical up to and including time  $t-1$ , then the corresponding decisions must be restricted to be equal at time  $t$  as we would not have any extra information to inform our decision. For the edge case  $t=1$ , we enforce that all decisions are the same, as no uncertainty has been revealed yet.

We provide a worked example of applying the multistage algorithm to the lot-sizing problem in Appendix C. This example demonstrates the substantial difference between our method and the method of Postek and Den Hertog (2014), which must employ various heuristics to avoid becoming overconservative.

**Figure 4.** (Color online) Partitioning of the uncertainty set in Example 2 using the multistage partitioning scheme.



*Notes.* At the second iteration, we have two distinct second-stage decisions, unlike in Figure 3 where we had only one. When we partition again, we grow to four distinct second-stage decisions, and nine distinct third-stage decisions.



## 5. Computational Experiments

The goal of this section is to demonstrate that, for problems of interest from the literature, the methods described in this paper are practical and valuable. To do so, we consider a two-stage binary capital budgeting problem (Section 5.1) and a multistage lot-sizing problem (Section 5.2). All methods were implemented in the Julia programming language using the JuMP (Lubin and Dunning 2015) and JuMPeR packages, with Gurobi 6.0.4 used as the solver (Gurobi Optimization 2015).

### 5.1. Capital Budgeting

Our first example is the two-stage capital budgeting problem described by Hanasusanto et al. (2014). We must decide which of  $N$  projects to pursue to maximize profits, but are constrained by a known budget  $B$ . The cost  $c_i$  and profit  $r_i$  for each project  $i$  are affine functions of uncertain parameters  $\xi$ ,

$$c_i(\xi) = (1 + \frac{1}{2} \Phi_i \cdot \xi) c_i^0 \quad \text{and} \quad r_i(\xi) = (1 + \frac{1}{2} \Psi_i \cdot \xi) r_i^0.$$

We have the option of pursuing a project either before or after these uncertain parameters are observed. If we pursue a project after they are observed, then we generate only a fraction  $\theta \in [0, 1)$  of the profit—a penalty for delaying the decision. This can be expressed as the adaptive binary optimization problem

$$\max_{z, \mathbf{x}} z$$

$$\text{subject to } \mathbf{r}(\xi) \cdot (\mathbf{x}^1 + \theta \mathbf{x}^2(\xi)) \geq z \quad \forall \xi \in \Xi$$

$$\mathbf{c}(\xi) \cdot (\mathbf{x}^1 + \mathbf{x}^2(\xi)) \leq B \quad \forall \xi \in \Xi$$

$$\mathbf{x}^1 + \mathbf{x}^2(\xi) \leq \mathbf{e} \quad \forall \xi \in \Xi$$

$$\mathbf{x}^1, \mathbf{x}^2(\xi) \in \{0, 1\}^N \quad \forall \xi \in \Xi,$$

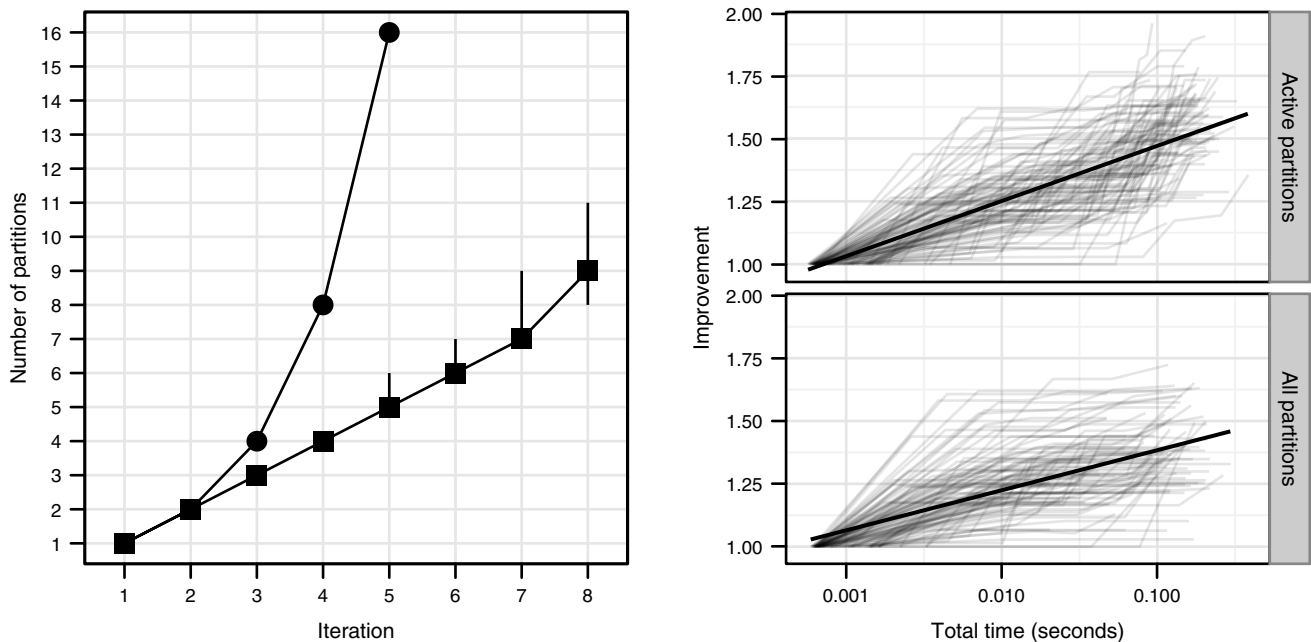
where  $\mathbf{x}^1$  is the first-stage decision to pursue now, and  $\mathbf{x}^2$  is the second-stage decision to pursue later. The uncertainty set is the hypercube  $\Xi = \{\xi \mid \xi \in [-1, 1]^4\}$ . We used the same distributions for the parameters as Hanasusanto et al. (2014), and set Gurobi's integrality gap tolerance to 1% for all instances. All RO problems were solved by reformulation to a deterministic problem.

We analyze our results from both the perspective of the decrease in gap versus total time, but also by “improvement” versus total time, where we use the same definition of improvement as Hanasusanto et al. (2014):

$$\frac{z_{\text{alg}}(\mathcal{T}^k) - z_{\text{alg}}(\mathcal{T}^1)}{z_{\text{alg}}(\mathcal{T}^1)}.$$

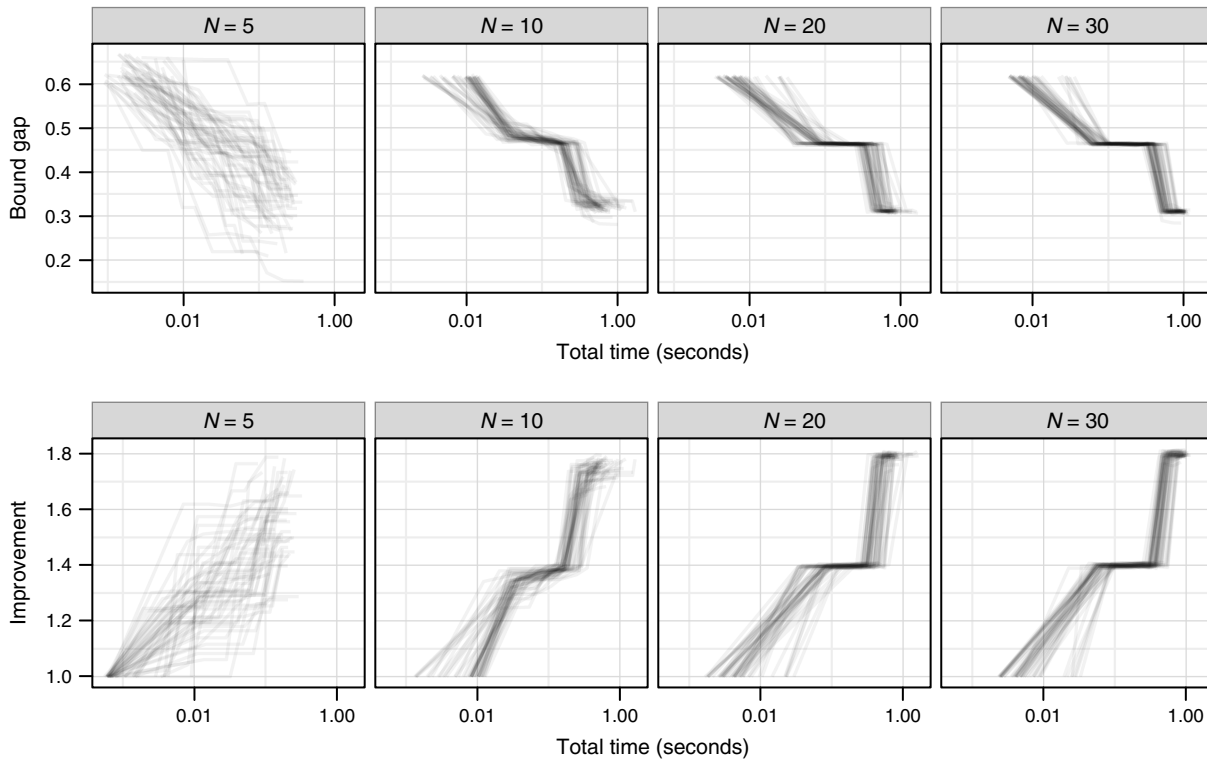
Before comparing results with Hanasusanto et al. (2014) for larger problem sizes, we first tried to understand the effects of only subpartitioning active partitions, which we define for this problem as the partitions that have an objective value that is within  $10^{-2}$  of the worst-case objective value. As can be seen in Figure 5, for 100 instances of size  $N = 5$ , we see that while subpartitioning all partitions leads to the number of partitions growing exponentially, focusing on

**Figure 5.** Performance of method when applied to the capital budgeting problem.



*Notes.* Left: median number of partitions versus iteration number if all partitions are subpartitioned (circles) or only active partitions are (squares). The top of the vertical line is the 90th percentile, and the bottom is the 10th percentile; there is no variation when all partitions are subpartitioned. Right: improvement versus time for the two approaches. Each semitransparent line represents the improvement across time for a single instance, and the solid black line is the line of best fit. Partitioning only on active partitions delivers better quality solutions, on average, than partitioning all partitions for the same computation time.

**Figure 6.** Improvement and bound gap vs. time for 50 instances of size  $N \in \{5, 10, 20, 30\}$ .



*Note.* Each instance is represented by a semitransparent line, and darker areas correspond to many overlapping lines.

only the active partitions leads to a near-linear increase in partition count. As discussed in Section 3.2, more important than partition count is the solution quality versus time. The right side of Figure 5 demonstrates this clearly, as four and eight iterations of each approach take approximately 0.1 seconds but the active partition-only approach delivers superior solutions on average.

For the purpose of comparison, we considered 50 random instances of each size  $N \in \{5, 10, 20, 30\}$ , and in Figure 6, we plot improvement and gap versus time for each instance. We found that there is little variability across instances in performance for  $N \geq 10$ , and that as  $N$  increases a pattern emerges where substantial initial improvement is made, before a multiple iteration stall, followed by further improvement. It seems reasonable to suggest that this pattern is driven by the ratio of  $N$  to the fixed dimension of the uncertainty set, which is why the behavior for  $N = 5$  differs substantially from the rest.

For instances of size  $N = 30$ , Hanasusanto et al. (2014) found an improvement of approximately 65% with their two-partition solution, 90% with their three-partition solution, and 110% with their four-partition solution. However, even with only  $N = 20$ , almost all instances failed to solve to optimality using their method within two hours. In contrast, we are able to consistently achieve an 80% improvement for instances of size  $N = 30$  in under one second (which corresponds to a 30% bound gap). We

suspect that this drastic difference in performance arises from the structure of the integer optimization problems being solved by the two methods. Our approach has a very simple block structure (one block per partition) not too dissimilar from the deterministic equivalent, whereas Hanasusanto et al. (2014) must introduce “big-M” constraints to allow for the partition structure to be included in the optimization, creating a computationally difficult problem. The two approaches are complementary, however, the method of Hanasusanto et al. (2014) would be preferable if having the best possible or a “simple” solution was important, such as when the second-stage decisions must be implemented as is and we cannot reoptimize. Additionally, while it was not the case in this problem, if there is only uncertainty in the objective function coefficients, then there are good theoretical reasons to suggest that the method of Hanasusanto et al. (2014) will perform better.

## 5.2. Multistage Lot Sizing

The multistage lot-sizing problem described in Bertsimas and Georghiou (2015) is a challenging AMIO problem that features a mix of continuous and binary decisions across multiple time stages. We have used smaller versions of this problem throughout this paper as a motivating example. The full formulation is provided in Bertsimas and Georghiou (2015), but we repeat the description here for

clarity. Our goal is to minimize the cost of stock ordered plus holding costs, subject to the need to satisfy uncertain demand. We must decide  $x^t$ , the (continuous, nonnegative) amount of stock ordered at time stage  $t$  for delivery at time stage  $t + 1$  at a unit cost of  $c_x$ , and  $y_n^t$ , a binary decision indicating whether to order a fixed amount  $q_n$  of stock for immediate delivery at time  $t$  at a unit cost of  $c_n$ . There are  $N$  different lot sizes to select from. We do not allow stock levels to become negative, and we impose a unit holding cost  $c_h$  for any stock remaining at the end of stage  $t$ . The uncertainty set is a box uncertainty set  $\Xi = \{\xi \mid \xi^1 = 1, l' \leq \xi^t \leq u^t \ \forall t \in \{2, \dots, T\}\}$ . For the continuous decisions  $\mathbf{x}$ , we start with an affine policy instead of a static policy, leading to a piecewise affine policy after partitioning (as described in Section 2.5). To aid understanding of the implementation of our method for this problem, we provide a worked example in Appendix C.

The problem size may grow rapidly, proportional to  $T^k$ , where  $k$  is the iteration count if our method is applied naively. We address this by only subpartitioning the active partitions. Additionally, many of the constraints only contain some of the uncertain parameters, and as a result, the set of active uncertain parameters is a set of infinite cardinality. To select a single active uncertain parameter per constraint, we randomly weight the components and solve a small optimization problem to find a random extreme point of the set. We again use reformulation to convert the RO problems to deterministic MIO problems. Gurobi's integrality gap tolerance was set to 1% for all instances, and the time limit was set to  $20T$  seconds for instances with  $T$  time stages.

We implemented our method and the method of Postek and Den Hertog (2014), and evaluated them over 50 randomly generated instances for  $T \in \{4, 6, 8, 10\}$  for  $N = 3$  (with results for  $N = 2$  being qualitatively similar). The method of Postek and Den Hertog (2014) has multiple heuristics and parameters that can be selected: we used the same that were used in that paper's computational experiments, with the exception of the lower-bound calculation where we used all found active uncertain parameters (which is what our method does). To make the running times of the two methods roughly equivalent, we used five iterations of the method of Postek and Den Hertog (2014), and three iterations of our method. We also sought to compare our results with Bertsimas and Georghiou (2015), who use a different definition of a bound gap than us:

$$\frac{z_{\text{alg}}(\mathcal{T}^k) - z_{\text{lower}}(\mathcal{T}^k)}{2(z_{\text{alg}}(\mathcal{T}^k) + z_{\text{lower}}(\mathcal{T}^k))}. \quad (30)$$

We have summarized the results in two ways. First, in Table 1, we compare our method with the methods of Postek and Den Hertog (2014) and Bertsimas and Georghiou (2015). Our method finds superior solutions in only a fraction of the time of Bertsimas and Georghiou (2015), suggesting that our method of approximating the

**Table 1.** Comparison of our method (with two or three iterations) vs. the methods of Postek and Den Hertog (2014) and Bertsimas and Georghiou (2015) for the multistage lot-sizing problem.

Method	$T$			
	4	6	8	10
Our method (2 iter.)				
Gap (%)	13.0	10.3	11.6	14.9
Time (s)	0.0	0.5	7.7	108.6
Our method (3 iter.)				
Gap (%)	11.4	9.3	11.3	14.9
Time (s)	0.2	2.0	52.4	309.3
Postek and Den Hertog (2014)				
Gap (%)	11.5	14.1	15.7	15.7
Time (s)	0.4	1.6	10.8	77.8
Bertsimas and Georghiou (2015)				
Gap (%)	17.2	34.5	37.6	—
Time (s)	3,381	9,181	28,743	—

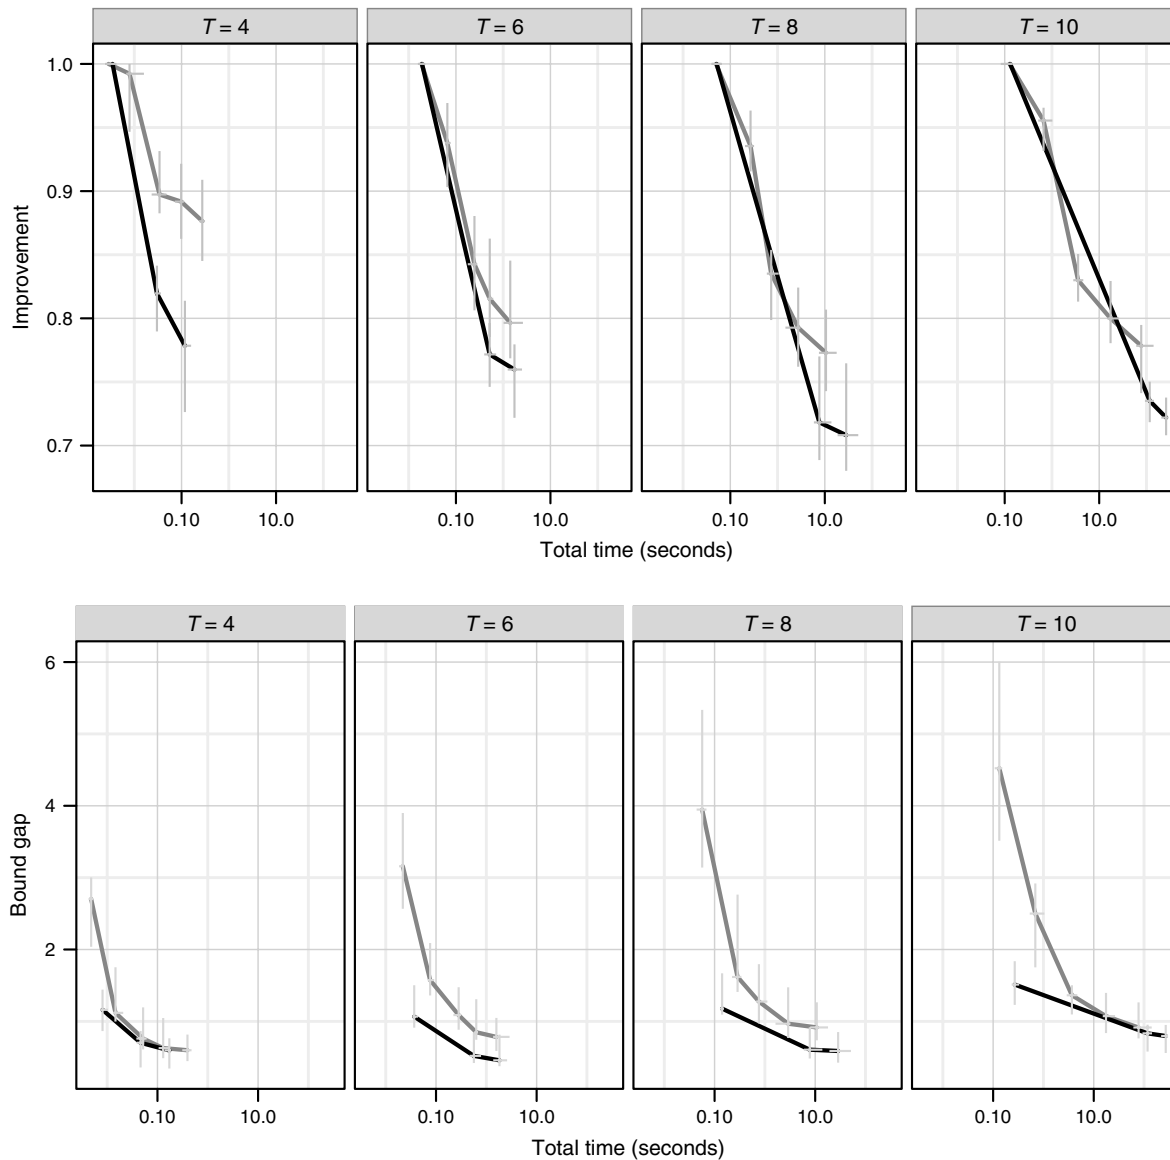
*Notes.* “Gap” is calculated as described in Bertsimas and Georghiou (2015), not as in this paper. Time for our method and Postek and Den Hertog (2014) is total time to solve all upper- and lower-bound problems up to and including that iteration; gap and time are medians for these two methods. Times for Bertsimas and Georghiou (2015) are obtained from a different computer but with roughly comparable performance, on instances from the same family.

optimal piecewise affine policy is more efficient than trying to directly optimize the “pieces.” Comparisons with Postek and Den Hertog (2014) are hard to make in this table format, because to be fair, we must compare total computation time versus solution quality, not iteration count versus solution quality. We visualize the full comparison, with quality measured by bound gap and improvement, in Figure 7. Here, we can see that the curves for our method, in general, lie close to or below the curves for Postek and Den Hertog (2014) across all instance sizes, suggesting that in this setting, our method will produce better solutions, on average, for the same amount of computational time invested.

## 6. Concluding Remarks

In this paper, we presented a novel partition-and-bound method for solving AMIO problems. We use finite partitioning to approximate the fully adaptable solution, and then use information from the solution of the finitely partitioned problem to further improve the partitioning and calculate lower bounds. The particular Voronoi diagram-like partitioning scheme we employ is inspired by the observation that, under some assumptions, a finite partitioning of the uncertainty set can only improve the solution quality if the partitioning separates regions of the uncertainty set that correspond to binding constraints. We find that our method delivers high-quality solutions for the time spent, even as the problem scales in number of time stages and

**Figure 7.** Time vs. improvement and bound gap for the multistage lot-sizing problem.



*Notes.* The black line and dark grey line correspond to our method and the method of Postek and Den Hertog (2014), respectively. All quantities are medians over 50 instances at each iteration, with the light grey error bars showing the 40th and 60th percentiles. The black curve is generally below the dark grey curve, indicating that our method obtains better results for the same amount of computational effort.

in the number of decision variables. The time per iteration does increase with the number of time stages, although we observed that the number of partitions created typically grows much slower than the worst case by focusing on active partitions. Our method should be viewed as one option among many for AMIO. For example if “good” solutions are required quickly, our method excels. If a near-fully adaptive solution is truly needed, other approaches may be more appropriate. We do note that, in a situation where we can reoptimize our decision at each time stage, perfectly identifying the fully adaptive solution may be unnecessary as the only important decision is the decision that must be made here and now.

Finally, we can draw analogies with the branch-and-bound technique for integer optimization. In branch-and-bound, we use the linear relaxation of an integer optimization problem to provide a lower bound on the true integer solution, use the information obtained from the solution to branch on fractional variables to obtain upper bounds, and repeat until we are satisfied with the bound gap. In integer optimization, the development of new cutting planes and presolve techniques has led to substantial improvements in the time to tighten these bounds gaps, and we believe there is exciting future work in developing new techniques for AMIO as extensions of this partition-and-bound approach.

## Acknowledgments

The authors thank the three anonymous reviewers for their helpful comments, Angelos Georgioudis for his feedback and encouragement, and Miles Lubin and Boris Braverman for their thoughts on partitioning schemes. Research was partially supported by the Office of Naval Research [Grant 021152-00001].

## Appendix A. Comparison of Partitioning Schemes for a Large Gap Example

Consider the following problem, which was presented in Bertsimas and Goyal (2010) to demonstrate large gaps between static and fully adaptive policies:

$$\begin{aligned} z(n) = \min_{\mathbf{x}^2 \geq 0, z} z \\ \text{subject to } \mathbf{e} \cdot \mathbf{x}^2(\xi) \leq z \quad \forall \xi \in \Xi \\ \mathbf{x}^2(\xi) \geq \xi \quad \forall \xi \in \Xi, \end{aligned}$$

where  $\Xi = \{\xi \mid \sum_{j=1}^n \xi_j \leq 1, \xi \geq 0\}$ . The initial static policy solution to this problem has objective value  $n$ , compared to the fully adaptive solution 1, and the active uncertain parameter for each component  $i$  of  $\mathbf{x}$  is the case where  $\xi_i = 1$  and all other components are zero.

**Our Method.** Our method creates  $n$  partitions, one for each component  $i$ . The partition for  $i = 1$  would be

$$\Xi(\hat{\xi}_1) = \left\{ \xi \mid \sum_{i=1}^n \xi_i \leq 1, \xi \geq 0, \xi_1 \geq \xi_2, \dots, \xi_1 \geq \xi_n \right\},$$

and similarly for the other partitions. If we consider the corresponding solution for this partition, we obtain  $\mathbf{x}_1^2 = (1, \frac{1}{2}, \dots, \frac{1}{2})$ , as while for the first component the worst case is unchanged ( $\xi_1 = 1$ ), the value of  $\xi_1$  must be at least as large as  $\xi_i$ , which limits the largest value any other component can take to  $\frac{1}{2}$ . Thus the problem at iteration 2 has  $n$  partitions and an objective of  $1 + (n-1)/2 = (n+1)/2$ , approximately half the initial objective for  $n \gg 1$ .

**Method of Postek and Den Hertog (2014).** The method of Postek and Den Hertog (2014) splits each partition into two subpartitions at each iteration, by selecting the two active uncertain parameters furthest apart. At the end of the first iteration, we can select the active uncertain parameters for components  $j = 1$  and  $j = 2$  without loss of generality. This results in two partitions:

$$\Xi(\hat{\xi}_1) = \left\{ \xi \mid \sum_{i=1}^n \xi_i \leq 1, \xi \geq 0, \xi_1 \geq \xi_2 \right\}$$

and

$$\Xi(\hat{\xi}_2) = \left\{ \xi \mid \sum_{i=1}^n \xi_i \leq 1, \xi \geq 0, \xi_1 \leq \xi_2 \right\}.$$

The corresponding solutions are  $\mathbf{x}_1^2 = (1, \frac{1}{2}, 1, \dots, 1)$  and  $\mathbf{x}_2^2 = (\frac{1}{2}, 1, \dots, 1)$ , respectively, giving an objective of  $n - \frac{1}{2}$ . At each subsequent iteration, another of the components will be restricted to be no more than  $\frac{1}{2}$ . Thus the objective at iteration  $k$  is  $n - (k-1)/2$  and the number of partitions is  $2^{k-1}$ . At iteration  $n$ , the objective will match that of our method  $n - (n-1)/2 = (n+1)/2$  and will have  $2^{n-1}$  partitions.

## Appendix B. Admissibility of Rule for Enforcing Nonanticipativity

We now demonstrate the rule described in Section 4.1 is sufficient to enforce nonanticipativity, but then show by example that it is overconservative: it will enforce nonanticipativity constraints needlessly in some cases.

**PROPOSITION 5.** If  $\mathbf{x}_i^t$  and  $\mathbf{x}_j^t$  are the decisions at time stage  $t \geq 2$  corresponding to the partitions  $\Xi(\hat{\xi}_i)$  and  $\Xi(\hat{\xi}_j)$ , respectively, then enforcing  $\mathbf{x}_i^t = \mathbf{x}_j^t$  (or not) according to the above decision rule ensures nonanticipativity.

**PROOF.** Consider the case where  $\hat{\xi}_i$  and  $\hat{\xi}_j$  are siblings.  $\mathbf{x}^t$  is a function of the uncertain parameters  $\xi^1, \dots, \xi^{t-1}$ . If  $t \leq t^{A,B}$  then  $\hat{\xi}_i^{1, \dots, t-1} = \hat{\xi}_j^{1, \dots, t-1}$ , and thus we cannot distinguish from the revealed uncertain parameters between  $\Xi(\hat{\xi}_i)$  and  $\Xi(\hat{\xi}_j)$ , and so must enforce the nonanticipativity constraints. If  $t > t^{A,B}$ , then a realization of the uncertain parameters  $\xi^1, \dots, \xi^{t-1}$  can be “assigned” to either one of the partitions (or neither) as a hyperplane based on the  $t^{A,B}$  component of  $\xi$  separates the two, and so we do not need to enforce nonanticipativity constraints.

If  $\hat{\xi}_i$  and  $\hat{\xi}_j$  are not siblings, we let  $\hat{\xi}_A \leftarrow \text{Parent}(\hat{\xi}_i)$  and  $\hat{\xi}_B \leftarrow \text{Parent}(\hat{\xi}_j)$ , and update  $t_{A,B}$ . If  $\hat{\xi}_A$  and  $\hat{\xi}_B$  are now siblings, then we know that all child partitions are separated by a constraint involving the terms  $\xi^{t^{A,B}}$ . If  $t \leq t^{A,B}$ , then using the same logic as above, we must conservatively add enforce the nonanticipativity constraints (even though they may be distinguished at a lower level). If  $t > t^{A,B}$ , then it can be assigned to one of the child partitions of either for the same reasons as above. If they are not siblings, we repeat and apply the same logic.  $\square$

As suggested in this proof, we may add constraints we do not need due to missing that two leafs might be separated in a way not easily determined just by the tree’s structure.

**EXAMPLE 3.** Consider the two-stage uncertainty set  $\Xi = \{(\xi^1, \xi^2) \mid 0 \leq \xi^1, \xi^2 \leq 10\}$ , with tree  $\mathcal{T}$

$$\mathcal{T} = \text{root} \begin{cases} (10, 0), & \begin{cases} (0, 0) & : A \\ (10, 0) & : B \end{cases} \\ (10, 10), & \begin{cases} (0, 10) & : C \\ (10, 10) & : D \end{cases} \end{cases}$$

where we have labelled each leaf A–D. The uncertainty sets corresponding to each leaf are

$$\begin{aligned} \Xi(\hat{\xi}_A) &= \{(\xi^1, \xi^2) \mid 0 \leq \xi^1 \leq 5, 0 \leq \xi^2 \leq 5\} \\ \Xi(\hat{\xi}_B) &= \{(\xi^1, \xi^2) \mid 5 \leq \xi^1 \leq 10, 0 \leq \xi^2 \leq 5\} \\ \Xi(\hat{\xi}_C) &= \{(\xi^1, \xi^2) \mid 0 \leq \xi^1 \leq 5, 5 \leq \xi^2 \leq 10\} \\ \Xi(\hat{\xi}_D) &= \{(\xi^1, \xi^2) \mid 5 \leq \xi^1 \leq 10, 5 \leq \xi^2 \leq 10\}, \end{aligned}$$

which we can manually inspect to determine that we must enforce the nonanticipativity constraints  $\mathbf{x}_A^2 = \mathbf{x}_C^2$  and  $\mathbf{x}_B^2 = \mathbf{x}_D^2$ . However, consider the application of the decision rule above to this problem for  $t = 2$ , that is,  $\mathbf{x}^2(\xi^1)$ .  $\hat{\xi}_A$  and  $\hat{\xi}_D$  are not siblings, so we must compare their parents (which are). The parents first differ in the second time stage, so we add a nonanticipativity constraint (as they appear indistinguishable knowing only  $\xi^1$ ). However, this constraint is not required, because only  $\xi^1 = 5$  could be in either  $\Xi(\hat{\xi}_A)$  or  $\Xi(\hat{\xi}_D)$ , but that is a boundary point. Thus the rule is overconservative in this case.  $\square$



## Appendix C. Comparison of Partitioning Schemes for Multistage Lot Sizing

The multistage lot-sizing problem of Bertsimas and Georgiou (2015) provides a good case study to understand the differences between the method in this paper and the proposal of Postek and Den Hertog (2014). We will walk through two iterations of both methods for  $T = 4$ , using the randomly generated uncertainty set  $\Xi = \{\xi^1 = 1, 21 \leq \xi^2 \leq 89, 13 \leq \xi^3 \leq 87, 4 \leq \xi^4 \leq 100\}$ . The other parameters are not directly relevant to the partitioning schemes, so we will not list them here.

We begin by solving the unpartitioned problem, which has an objective value of  $1,830\frac{2}{3}$ . The unique active uncertain parameters we extract are

$$\hat{\xi}_1 = (1, 21, 13, 4), \quad \hat{\xi}_3 = (1, 89, 13, 4), \quad \text{and} \quad \hat{\xi}_5 = (1, 89, 87, 100). \\ \hat{\xi}_2 = (1, 21, 87, 4), \quad \hat{\xi}_4 = (1, 89, 87, 4).$$

At this point, the two methods diverge, as our method creates five partitions while the other creates two.

### Our method

**Iteration 1.** Our method considers each pair of active uncertain parameters, pairwise, when constructing partitions. We now detail the construction for  $\hat{\xi}_3$ . Consider first  $\hat{\xi}_3$  and  $\hat{\xi}_1$ :  $t_{3,1}$ , the first time stage they differ, is 2. Thus the first constraint is  $\xi^2 \geq 55$ , with the same constraint produced for  $\hat{\xi}_2$  as well. For  $\hat{\xi}_4$  and  $\hat{\xi}_5$ , we find  $t_{3,4} = t_{3,5} = 3$ , giving the constraint  $\xi^3 \leq 50$ . We can now express all five partitions:

$$\Xi(\hat{\xi}_1) = \{\xi^1 = 1, 21 \leq \xi^2 \leq 55, 13 \leq \xi^3 \leq 50, 4 \leq \xi^4 \leq 100\}, \\ \Xi(\hat{\xi}_2) = \{\xi^1 = 1, 21 \leq \xi^2 \leq 55, 50 \leq \xi^3 \leq 87, 4 \leq \xi^4 \leq 100\}, \\ \Xi(\hat{\xi}_3) = \{\xi^1 = 1, 55 \leq \xi^2 \leq 89, 13 \leq \xi^3 \leq 50, 4 \leq \xi^4 \leq 100\}, \\ \Xi(\hat{\xi}_4) = \{\xi^1 = 1, 55 \leq \xi^2 \leq 89, 50 \leq \xi^3 \leq 87, 4 \leq \xi^4 \leq 52\}, \\ \Xi(\hat{\xi}_5) = \{\xi^1 = 1, 55 \leq \xi^2 \leq 89, 50 \leq \xi^3 \leq 87, 52 \leq \xi^4 \leq 100\}.$$

The final consideration is nonanticipativity constraints. We must enforce the following constraints, all of which are relatively intuitive as all our constraints include only a single time stage:  $x_1^2 = x_2^2$ ,  $x_3^3 = x_4^3 = x_5^3$ ,  $x_4^4 = x_5^4$ , with similar constraints on  $y$ . Naturally, the decision  $x^1$  is the same for all partitions as well.

**Iteration 2.** We improve the objective from  $1,830\frac{2}{3}$  to  $1,672$  thanks to our partitioning scheme. The “contributions” to the objective from each partition are  $1,426\frac{1}{3}$ ,  $1,672$ ,  $1,627$ ,  $1,530\frac{2}{3}$ , and  $1,672$ , respectively. Only two of these partitions are “active partitions,” so we will further divide only these two. Here, we consider partition 2 only for the sake of exposition. This partition has the active uncertain parameters  $\hat{\xi}_{2,1} = (1, 21, 87, 4)$ ,  $\hat{\xi}_{2,2} = (1, 21, 50, 4)$ ,  $\hat{\xi}_{2,3} = (1, 55, 87, 4)$ , and  $\hat{\xi}_{2,4} = (1, 21, 50, 100)$ . These produce the following partitions, all of which are subpartitions of the initial partition:

$$\Xi(\hat{\xi}_{2,1}) = \{\xi^1 = 1, 21 \leq \xi^2 \leq 38, 68.5 \leq \xi^3 \leq 87, 4 \leq \xi^4 \leq 100\}, \\ \Xi(\hat{\xi}_{2,2}) = \{\xi^1 = 1, 21 \leq \xi^2 \leq 38, 50 \leq \xi^3 \leq 68.5, 4 \leq \xi^4 \leq 52\}, \\ \Xi(\hat{\xi}_{2,3}) = \{\xi^1 = 1, 38 \leq \xi^2 \leq 55, 50 \leq \xi^3 \leq 87, 4 \leq \xi^4 \leq 100\}, \\ \Xi(\hat{\xi}_{2,4}) = \{\xi^1 = 1, 21 \leq \xi^2 \leq 38, 50 \leq \xi^3 \leq 68.5, 52 \leq \xi^4 \leq 100\}.$$

For nonanticipativity, we must first enforce the constraints of the first iteration. As we are not further partitioning partition 1 at this time, this translates to  $x_1^2 = x_{2,1}^2 = x_{2,2}^2 = x_{2,3}^2 = x_{2,4}^2$ . We must also add new nonanticipativity constraints, namely,  $x_{2,1}^3 = x_{2,2}^3 = x_{2,4}^3$  and  $x_{2,2}^3 = x_{2,3}^3$ .

### Method of Postek and Den Hertog (2014)

**Iteration 1.** The first step is to identify the  $t$  we will be using in the split. As described in Section 5.1 of Postek and Den Hertog (2014), we aim to find the component of the uncertain parameters that has maximal “dispersion” for a  $t$  such that  $t_{\max} \leq t \leq t_{\max} + q$ . The initial  $t_{\max}$  for the set is 0, and the authors used  $q = 2$  in their computational experiments. We are left with just  $t = 2$  as the uncertainty set is fixed for  $t = 1$ .

We now proceed to creating a “2-SH.” Using “Heuristic 1” of Section 5.2, we select  $\hat{\xi}_2$  and  $\hat{\xi}_5$  as having maximal pairwise distance over the selected time components, giving the hyperplane  $\xi^2 = 55$ . Each partition has  $t_{\max} = 2$ , and no nonanticipativity constraints are required apart from the trivial one on  $x^1$ .

**Iteration 2.** The objective value remains at  $1,830\frac{2}{3}$ , with both partitions individually also “contributing” this objective. We will focus on the partition that was defined by the constraint  $\xi^2 \geq 55$ .

This partition has the active uncertain parameters  $\hat{\xi}_{2,1} = (1, 55, 13, 4)$ ,  $\hat{\xi}_{2,2} = (1, 89, 13, 4)$ ,  $\hat{\xi}_{2,3} = (1, 89, 87, 4)$ , and  $\hat{\xi}_{2,4} = (1, 89, 13, 100)$ . We now search for the  $t$  to partition on between  $t_{\max} = 2$  and  $t_{\max} + q = 4$ . The maximum “dispersion” occurs for  $t = 4$ , so we will construct a “4-SH.” The two  $\hat{\xi}$  with maximum pairwise distance are  $\hat{\xi}_{2,3}$  and  $\hat{\xi}_{2,4}$ , so the hyperplane we construct is  $-74\xi^3 + 96\xi^4 = 1,292$ . We would thus need  $\xi^4$  to distinguish which partition we are in, so we need nonanticipativity constraints  $x_{2,3}^2 = x_{2,4}^2$  and  $x_{2,3}^3 = x_{2,4}^3$ .

## References

- Aurenhammer F (1991) Voronoi diagrams—A survey of a fundamental geometric data structure. *ACM Comput. Surveys (CSUR)* 23(3): 345–405.
- Ben-Tal A, El Ghaoui L, Nemirovski A (2009) *Robust Optimization* (Princeton University Press, Princeton, NJ).
- Ben-Tal A, Golany B, Nemirovski A, Vial J-P (2005) Retailer-supplier flexible commitments contracts: A robust optimization approach. *Manufacturing Service Oper. Management* 7(3):248–271.
- Ben-Tal A, Goryashko A, Guslitzer E, Nemirovski A (2004) Adjustable robust solutions of uncertain linear programs. *Math. Programming* 99(2):351–376.
- Bertsimas D, Bidkhori H (2015) On the performance of affine policies for two-stage adaptive optimization: A geometric perspective. *Math. Programming* 153(2):577–594.
- Bertsimas D, Caramanis C (2007) Adaptability via sampling. *Proc. 46th IEEE Conf. Decision and Control* (IEEE, Piscataway, NJ), 4717–4722.
- Bertsimas D, Caramanis C (2010) Finite adaptability in multistage linear optimization. *IEEE Trans. Automatic Control* 55(12):2751–2766.
- Bertsimas D, Georgiou A (2014) Binary decision rules for multistage adaptive mixed-integer optimization. *Optim. Online*.
- Bertsimas D, Georgiou A (2015) Design of near optimal decision rules in multistage adaptive mixed-integer optimization. *Oper. Res.* 63(3):610–627.
- Bertsimas D, Goyal V (2010) On the power of robust solutions in two-stage stochastic and adaptive optimization problems. *Math. Oper. Res.* 35(2):284–305.

- Bertsimas D, Goyal V (2012) On the power and limitations of affine policies in two-stage adaptive optimization. *Math. Programming* 134(2): 491–531.
- Bertsimas D, Brown DB, Caramanis C (2011a) Theory and applications of robust optimization. *SIAM Rev.* 53(3):464–501.
- Bertsimas D, Dunning I, Lubin M (2015) Reformulation versus cutting-planes for robust optimization. *Computational Management Sci.* 13(2):195–217.
- Bertsimas D, Iancu D, Parrilo P (2010) Optimality of affine policies in multistage robust optimization. *Math. Oper. Res.* 35(2):363–394.
- Bertsimas D, Iancu D, Parrilo P (2011b) A hierarchy of near-optimal policies for multistage adaptive optimization. *IEEE Trans. Automatic Control* 56(12):2809–2824.
- Bertsimas D, Litvinov E, Sun XA, Zhao J, Zheng T (2013) Adaptive robust optimization for the security constrained unit commitment problem. *IEEE Trans. Power Systems* 28(1):52–63.
- Chen X, Sim M, Sun P (2007) A robust optimization perspective on stochastic programming. *Oper. Res.* 55(6):1058–1071.
- Chen X, Sim M, Sun P, Zhang J (2008) A linear decision-based approximation approach to stochastic programming. *Oper. Res.* 56(2):344–357.
- Fischetti M, Monaci M (2012) Cutting plane versus compact formulations for uncertain (integer) linear programs. *Math. Programming Comput.* 4(3):239–273.
- Goulart PJ, Kerrigan EC, Maciejowski JM (2006) Optimization over state feedback policies for robust control with constraints. *Automatica* 42(4):523–533.
- Gurobi Optimization, Inc. (2015) Gurobi optimizer reference manual. <http://www.gurobi.com>.
- Hadjiyiannis MJ, Goulart PJ, Kuhn D (2011) A scenario approach for estimating the suboptimality of linear decision rules in two-stage robust optimization. *50th IEEE Conf. Decision and Control Eur. Control Conf., CDC-ECC*, (IEEE, Piscataway, NJ), 7386–7391.
- Hanasusanto GA, Kuhn D, Wiesemann W (2014) Two-stage robust integer programming. *Optim. Online*.
- Lee D-T, Drysdale RL III (1981) Generalization of Voronoi diagrams in the plane. *SIAM J. Comput.* 10(1):73–87.
- Lubin M, Dunning I (2015) Computing in operations research using Julia. *INFORMS J. Comput.* 27(2):238–248.
- Postek K, Den Hertog D (2014) Multi-stage adjustable robust mixed-integer optimization via iterative splitting of the uncertainty set. *Center Discussion Paper Series No.* 2016-006.
- Song Y, Luedtke J (2014) An adaptive partition-based approach for solving two-stage stochastic programs with fixed recourse *SIAM J. Optim.* 25(3):1344–1367.
- Vayanos P, Kuhn D, Rustem B (2011) Decision rules for information discovery in multi-stage stochastic programming. *50th IEEE Conf. Decision and Control Eur. Control Conf., CDC-ECC*, (IEEE, Piscataway, NJ), 7368–7373.
- Zeng B, Zhao L (2013) Solving two-stage robust optimization problems using a column-and-constraint generation method. *Oper. Res. Lett.* 41(5):457–461.

**Dimitris Bertsimas** is the Boeing Professor of Operations Research and the codirector of the Operations Research Center at Massachusetts Institute of Technology. He is a member of the National Academy of Engineering, and the recipient of numerous research and teaching awards. The present paper is part of his ongoing research on robust optimization.

**Iain Dunning** is a doctoral candidate at the Massachusetts Institute of Technology Operations Research Center. This paper is part of his doctoral thesis on new methods for robust and adaptive optimization.